



# AI 3-in-1: Agents, RAG, and Local Models



Presented by Brent Laster &  
Tech Skills Transformations LLC


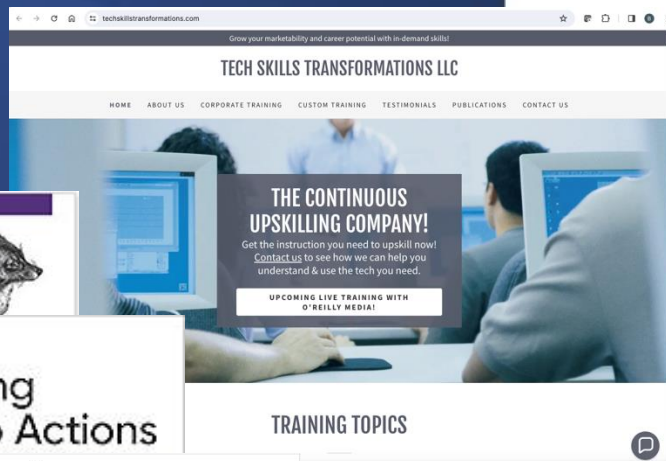
© 2025 Brent C. Laster & Tech Skills Transformations LLC



All rights reserved



# About me



- Founder, Tech Skills Transformations LLC
- <https://getskillsnow.com>
- [info@getskillsnow.com](mailto:info@getskillsnow.com)
- Long career in corporate as dev, manager, and director in DevOps and other areas


## Author

- O'Reilly "reports"
- Books
  - Professional Git
  - Jenkins 2 - Up and Running
  - Learning GitHub Actions
  - Learning GitHub Copilot
  - AI-Enabled SDLC

## Speaker


- Social media

- ☐ **LinkedIn:** [brentlaster](#)
- ☐ **X:** [@BrentCLaster](#)
- ☐ **Bluesky:** [brentclaster.bsky.social](#)
- ☐ **GitHub:** [brentlaster](#)



LISTED BELOW ARE A FEW OF THE TECHNOLOGIES FOR WHICH WE OFFER HANDS-ON TRAINING. THESE CAN BE CUSTOMIZED FOR ANY SIZE TEAM FROM 1-100 AND FOR ANY LEVEL FROM BEGINNER TO ADVANCED.

- ArgoCD
- Containers
- Docker
- Gerrit
- Git
- GitHub Actions
- GitHub Codespaces
- GitHub Copilot
- GitHub Foundations
- GitHub Security
- GitLab
- GitOps
- Gradle
- Grafana
- Helm
- Jenkins
- Kubernetes
- Kustomize
- LLMs
- Prometheus
- Tekton
- VS Code





## Running models locally





# Why run models locally?

5

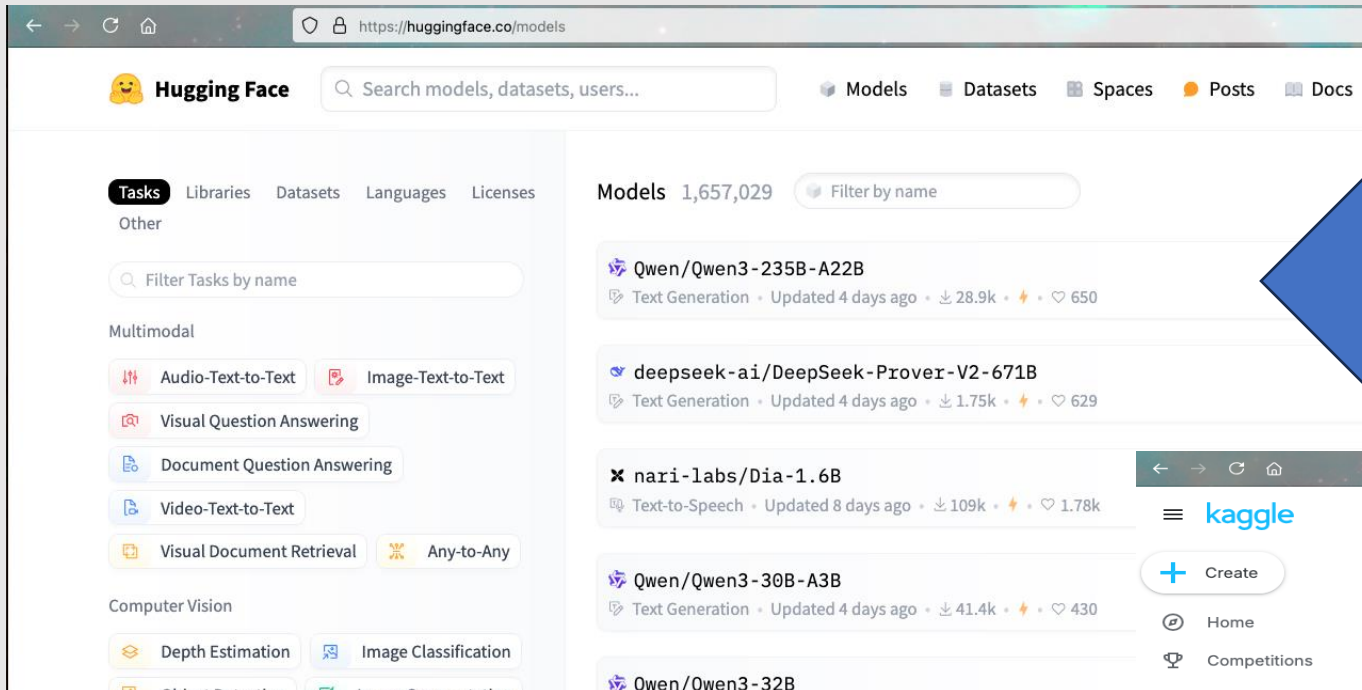
- Privacy - no need to share data
- Gives you control over setup, configuration, and customization options
  - Can tailor LLM to your needs, experiment with settings, integrate into your infra
- Can easily swap between different models for different tasks
- Work in offline mode
- Cost savings
  - No charges for subscriptions or API calls
- No censoring of results



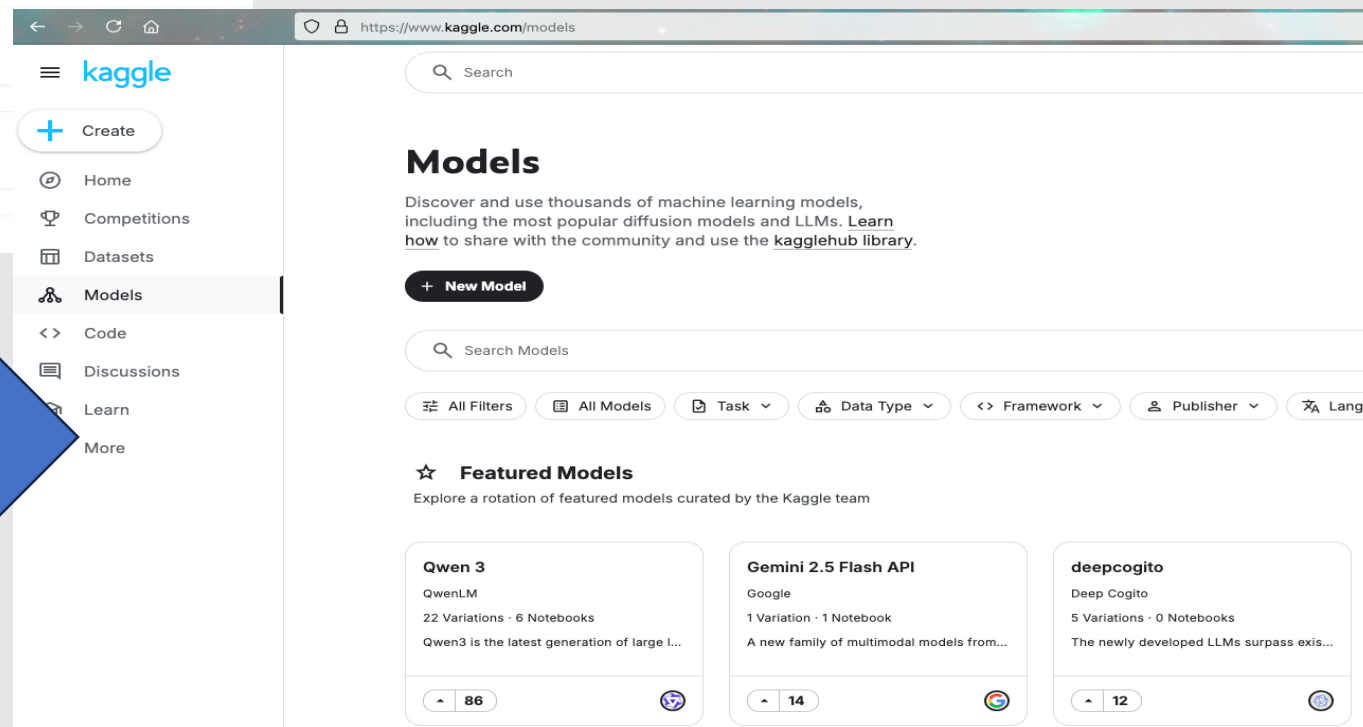


# Where to get models +

6



<http://huggingface.co/models>



<http://kaggle.com/models>







# Options for running LLMs locally

7

- GPT4All - <https://github.com/nomic-ai/gpt4all>
- LM Studio - <https://lmstudio.ai>
- Jan AI - <https://jan.ai>
- llama.cpp - <https://github.com/ggerganov/llama.cpp>
- LlamaFile - <https://github.com/Mozilla-Ocho/llamafile>
- Ollama - <https://ollama.com/>
- HuggingFace Transformers - <https://huggingface.co/docs/transformers>
- More!



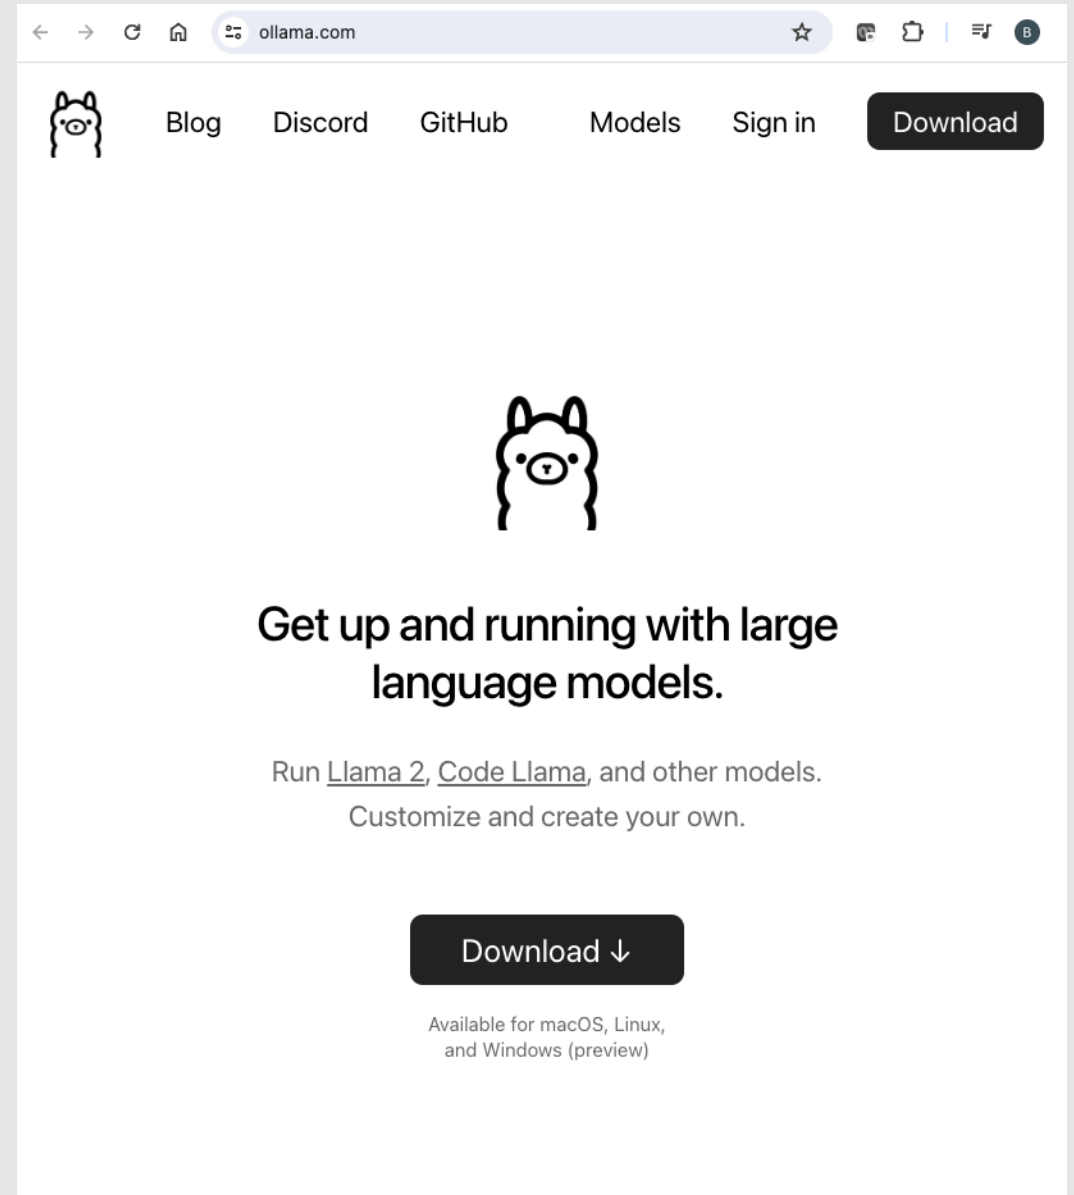
**LlamaFile**  
Local Executable GPT



**Transformers**



- Command line tool for downloading, exploring and using LLMs on local machine
- open source
- supports most of Hugging Face's popular models
- allows uploading new ones
- Links:
  - main site: <https://ollama.com>
  - GitHub: <https://github.com/ollama/>
- Advantages
  - speeds up and simplifies
    - » model selection and download
    - » configuring endpoints
    - » integration with Python or JavaScript codebase


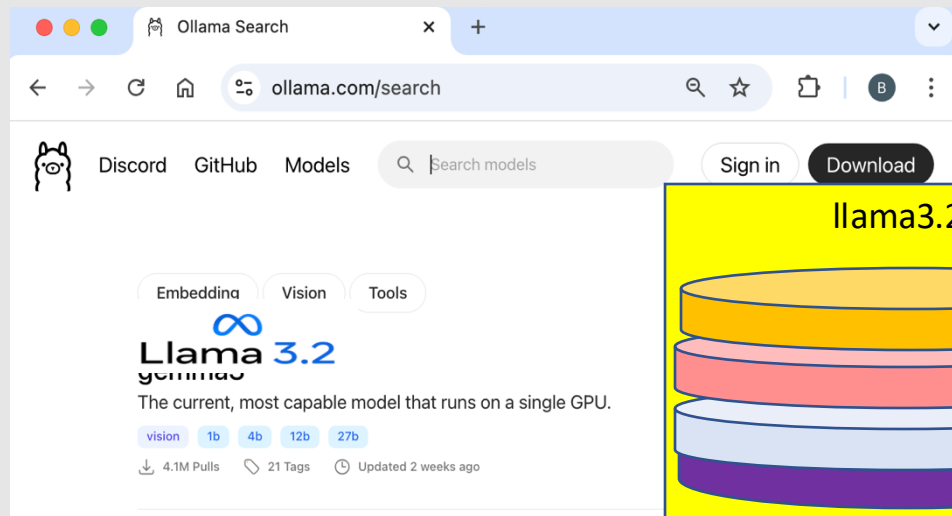







# Working with Ollama #1

9



*ollama pull*







# Working with Ollama #2

10


*>>> Briefly explain what an AI model is*

```
PROBLEMS  TERMINAL  ...
bash - 3in1 + - [ ] [X] ... ^ X
(py_env) @gwstudent2 → /workspaces/3in1 (main) $
```

*ollama run*



*>>> query*





# Working with Ollama #3

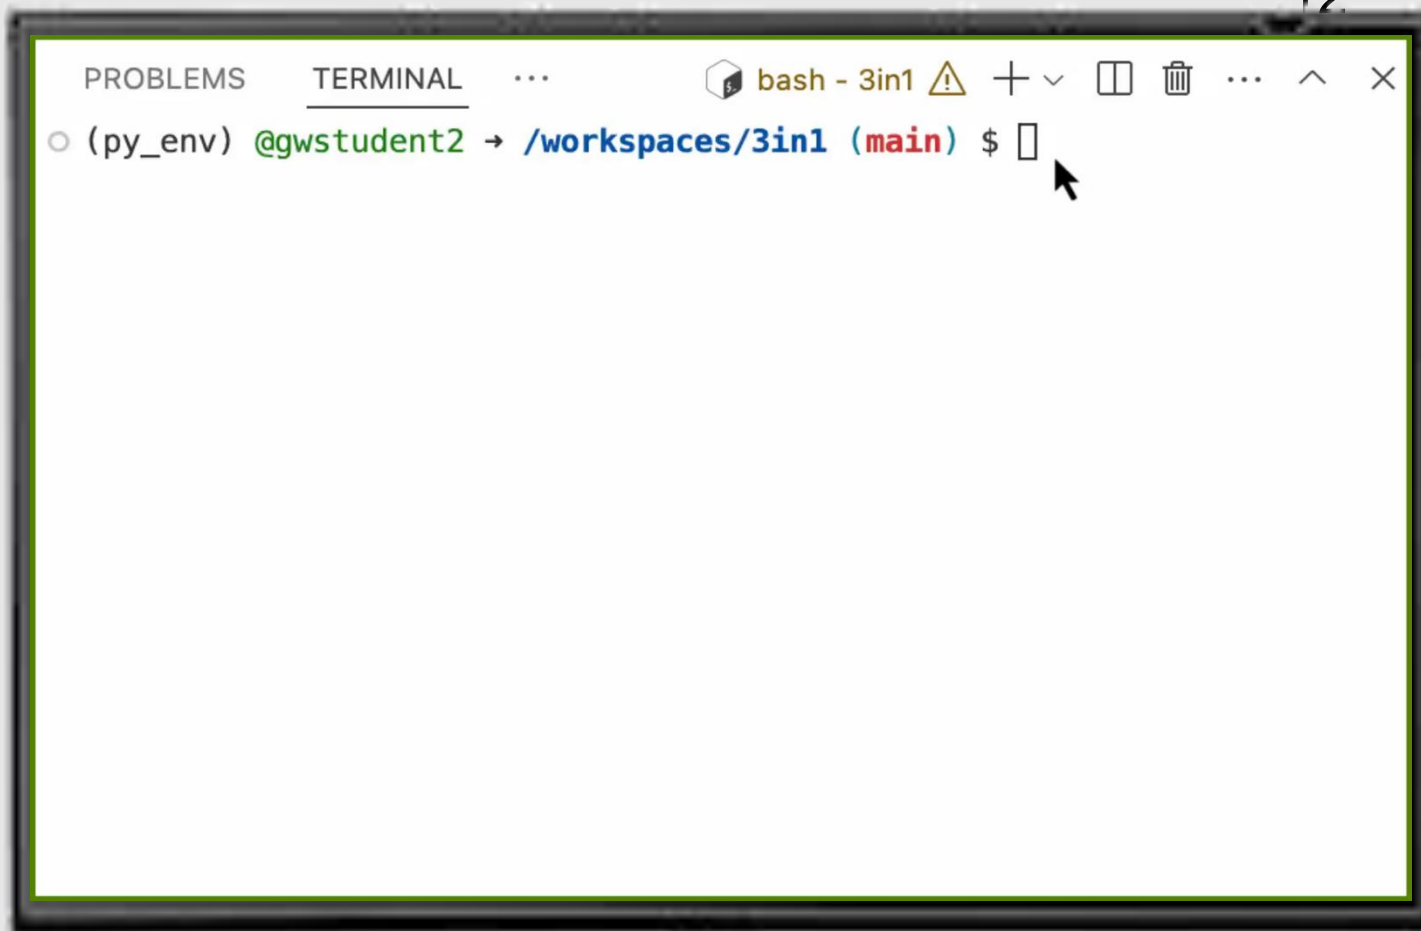
12

```
from openai import OpenAI

# Initialize the client to talk to your local Ollama server
client = OpenAI(
    base_url="http://localhost:11434/v1",
    api_key="n-a"
)

# Send a chat completion request
completion = client.chat.completions.create(
    model="llama3.2",
    messages=[
        {"role": "system", "content": "Always answer in 3 bullet points."},
        {"role": "user", "content": "Tell me what AI is."}
    ],
    temperature=0.7,
)

# Print out the assistant's reply
print(completion.choices[0].message.content)
```



*ollama serve*





## Demo #1 – Simple program to work with local model





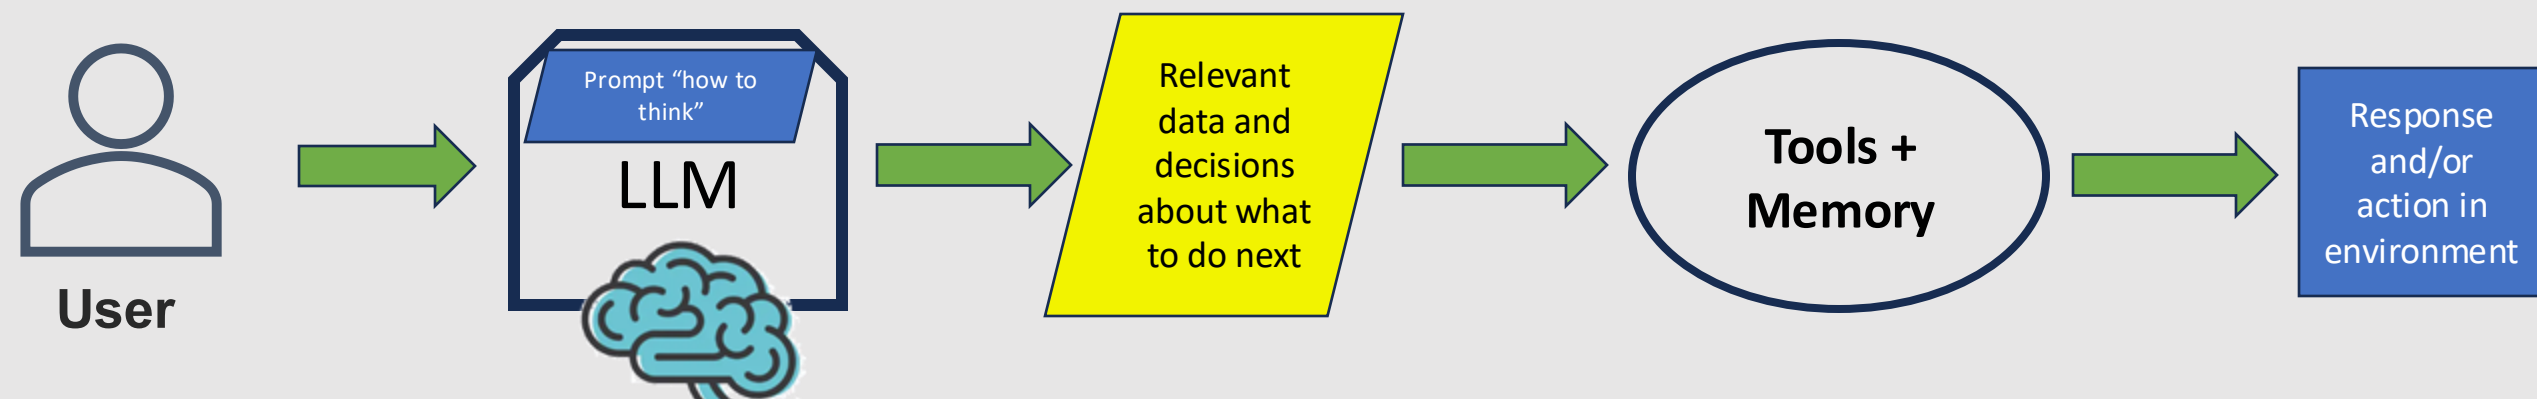
# Agents





# What is an AI Agent?

- A **system** that operates within an **environment** by using **sensors** to **perceive** information, a **decision-making mechanism** to **process and reason about the data**, and **actuators** to **take actions that influence or update/respond to the environment**
- This interaction enables the agent to achieve specific goals autonomously while continuously learning and adapting over time
- Agents use LLMs to identify key data, drive decisions, and communicate naturally





# Architectural Features of AI Agents

## Planning



- AI autonomously outlines and executes a logical series of steps for accomplishing a given objective.
- Provides the AI with a way to dynamically adapt its approach based on real-time data and feedback..
- Might employ reflection to evaluate and improve responses
- **Example:** A research agent plans search → summarize → generate report.

## Tool Use



- AI agents interact with external APIs, databases, and functions.
- Enhances LLMs by providing access to real-world knowledge.
- Reduces hallucinations by using retrieval-augmented generation (RAG).
- **Example:** Calling a Python function to perform complex calculations.

## Memory



- Short-term handles tasks; long term stores knowledge and experience
- Memory ensures consistency and efficiency in multi-step decisions
- Memory recalls preferences to enhance personalization and user experience
- **Example:** Storing user preferences for future reference or personalized responses



**system\_message=“””You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.**

**You have access to the following tools:**

**Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string**

**You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.**

**You should first reflect with “Thought: {your\_thoughts}” on the current query, then (if necessary), call a tool with the proper JSON formatting “Action: {JSON\_BLOB}”, or else print your final answer starting with the prefix “Final Answer:”“””**







# Agent Example



**system\_message**="""You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses. You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" ""

18



User

What's the weather in Paris?



LLM

## Chain of Thought – Step 1: Interpret User Query

Thought: "The user is asking about the weather in Paris. I need to extract 'Paris' as the location.

Action: Extracted location = "Paris"



Weather  
Search Tool

AI Agent





# Agent Example


**system\_message="""**You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.  
You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" ""

19



User

What's the weather in Paris?



LLM

**Chain of Thought – Step 2: Decide to use tool**  
Thought: "I need real-time data, so I will call the 'find\_weather' tool. First, I need to get the latitude and longitude for the tool call."

```
AIResponse(  
  tool_calls=[  
    {  
      name:  
        "find_weather"  
      parameters: {  
        latitude:  
          "48.8566",  
        longitude:  
          "2.3522",  
      },  
      id: "call_tool123",  
      type: "tool_invoke"  
    }  
  ]  
)
```

**Weather  
Search Tool**

AI Agent





# Agent Example


**system\_message="""**You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.  
You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" ""

20



User

What's the weather in Paris?



LLM

```
AIResponse(  
  tool_calls=[  
    name:  
    "find_weather"  
    parameters: {  
      latitude:  
      "48.8566",  
      longitude:  
      "2.3522",  
    },  
    id: "call_tool123",  
    type: "tool_invoke"  
  ]  
)
```



Weather  
Search Tool

Agent parses LLM output  
identifies JSON tool call,  
parses it, forms it into  
actual tool call

```
{  
  name:  
  "find_weather"  
  parameters: {  
    latitude:  
    "48.8566",  
    longitude:  
    "2.3522",  
  },  
  id: "call_tool123",  
  type: "tool_invoke"  
}
```

AI Agent





# Agent Example


**system\_message="""**You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.  
You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" **""**

21



User

What's the weather in Paris?



LLM

```
AIResponse(  
  tool_calls=[  
    {  
      name:  
        "find_weather"  
      parameters: {  
        latitude:  
          "48.8566",  
        longitude:  
          "2.3522",  
      },  
      id: "call_tool123"  
    }  
  ]  
)
```

Agent executes tool call



Weather  
Search Tool

```
{  
  name:  
    "find_weather"  
  parameters: {  
    latitude:  
      "48.8566",  
    longitude:  
      "2.3522",  
  },  
  id: "call_tool123",  
  type: "tool_invoke"  
}
```

AI Agent





# Agent Example


**system\_message**="""You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses. You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" ""

22



User

What's the weather in Paris?



LLM

```
AIResponse(  
  tool_calls=[  
    {  
      name:  
        "find_weather"  
      parameters: {  
        latitude:  
          "48.8566",  
        longitude:  
          "2.3522",  
      },  
      id: "call_tool123",  
      type: "tool_invoke"  
    }  
  ]  
)
```

```
{  
  name:  
    "find_weather"  
  parameters: {  
    latitude:  
      "48.8566",  
    longitude:  
      "2.3522",  
  },  
  id: "call_tool123",  
  type: "tool_invoke"  
}
```

```
ToolResponse(  
  content="53 and  
  rainy",  
  name="find_weather",  
  tool_invoke_id:  
    "call_tool123"
```

Weather tool returns result



Weather  
Search Tool

AI Agent





# Agent Example


**system\_message="""**You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.  
You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" **""**

23



User

What's the weather in Paris?



LLM

```
AIResponse(  
  tool_calls=[  
    {  
      name:  
        "find_weather"  
      parameters: {  
        latitude:  
          "48.8566",  
        longitude:  
          "2.3522",  
      },  
      id: "call_tool123",  
      type: "tool_invoke"  
    }  
  ]  
)
```

Agent includes tool output in message/prompt back to model

```
ToolResponse(  
  content="53 and rainy",  
  name="find_weather",  
  tool_invoke_id:  
    "call_tool123"
```



Weather  
Search Tool

```
{  
  name:  
    "find_weather"  
  parameters: {  
    latitude:  
      "48.8566",  
    longitude:  
      "2.3522",  
  },  
  id: "call_tool123",  
  type: "tool_invoke"
```

AI Agent





# Agent Example


24

**system\_message="""**You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.  
You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" ""



User

What's the weather in Paris?



LLM

```
AIResponse(  
  tool_calls=[  
    {  
      name:  
        "find_weather"  
      parameters: {  
        latitude:  
          "48.8566",  
        longitude:  
          "2.3522",  
      },  
      id: "call_tool123",  
      type: "tool_invoke"  
    }  
  ]  
)
```

## Chain of Thought – Step 3 : Interpret JSON Response

Thought: "The tool returned weather data for Paris. I will summarize the information concisely."

```
name="find_weather",  
tool_invoke_id:  
"call_tool123"
```



Weather  
Search Tool

AI Agent

```
{  
  name:  
    "find_weather"  
  parameters: {  
    latitude:  
      "48.8566",  
    longitude:  
      "2.3522",  
  },  
  id: "call_tool123",  
  type: "tool_invoke"  
}
```







# Agent Example


25

**system\_message="""**You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.  
You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" **""**

What's the weather in Paris?



User



```
AIResponse(  
  tool_calls=[  
    {  
      name:  
        "find_weather"  
      parameters: {  
        latitude:  
          "48.8566",  
        longitude:  
          "2.3522",  
      },  
      id: "call_tool123",  
      type: "tool_invoke"  
    }  
  ]  
)
```

```
ToolResponse(  
  content="53 and  
  rainy",  
  name="find_weather",  
  tool_invoke_id:  
    "call_tool123"
```



```
{  
  name:  
    "find_weather"  
  parameters: {  
    latitude:  
      "48.8566",  
    longitude:  
      "2.3522",  
  },  
  id: "call_tool123",  
  type: "tool_invoke"
```

```
AIFinalResponse(  
  content="The  
  current weather in Paris  
  is 53 degrees Celsius  
  with light rain."  
)
```

AI Agent





# Agent Example


26

**system\_message="""**You are an AI assistant designed to help users accurately and efficiently. Your primary goal is to provide precise, helpful, and clear responses.  
You have access to the following tools:  
Tool Name: find\_weather, Description: Get weather for a location., Arguments: latitude: float, longitude: float, Outputs: string  
You should think step by step in order to fulfill the objective with a reasoning process divided into Thought/Action/Observation. This cycle can repeat multiple times if needed.  
You should first reflect with "Thought: {your\_thoughts}" on the current query, then (if necessary), call a tool with the proper JSON formatting "Action: {JSON\_BLOB}", or else print your final answer starting with the prefix "Final Answer:" **""**

What's the weather in Paris?



User



```
AIResponse(  
  tool_calls=[  
    {  
      name:  
        "find_weather"  
      parameters: {  
        location: "Paris",  
      },  
      id: "call_tool123",  
      type: "tool_invoke"  
    }  
  ]  
)
```

```
ToolResponse(  
  content="53 and  
  rainy",  
  name="find_weather",  
  tool_invoke_id:  
    "call_tool123"
```

Weather  
Search Tool

```
{  
  name:  
    "find_weather"  
  parameters: {  
    latitude:  
      "48.8566",  
    longitude:  
      "2.3522",  
  },  
  id: "call_tool123",  
  type: "tool_invoke"
```

```
AIFinalResponse(  
  content="The  
  current weather in Paris  
  is 53 degrees Celsius  
  with light rain."  
)
```

AI Agent





## Demo #2 – Adding agency to our code





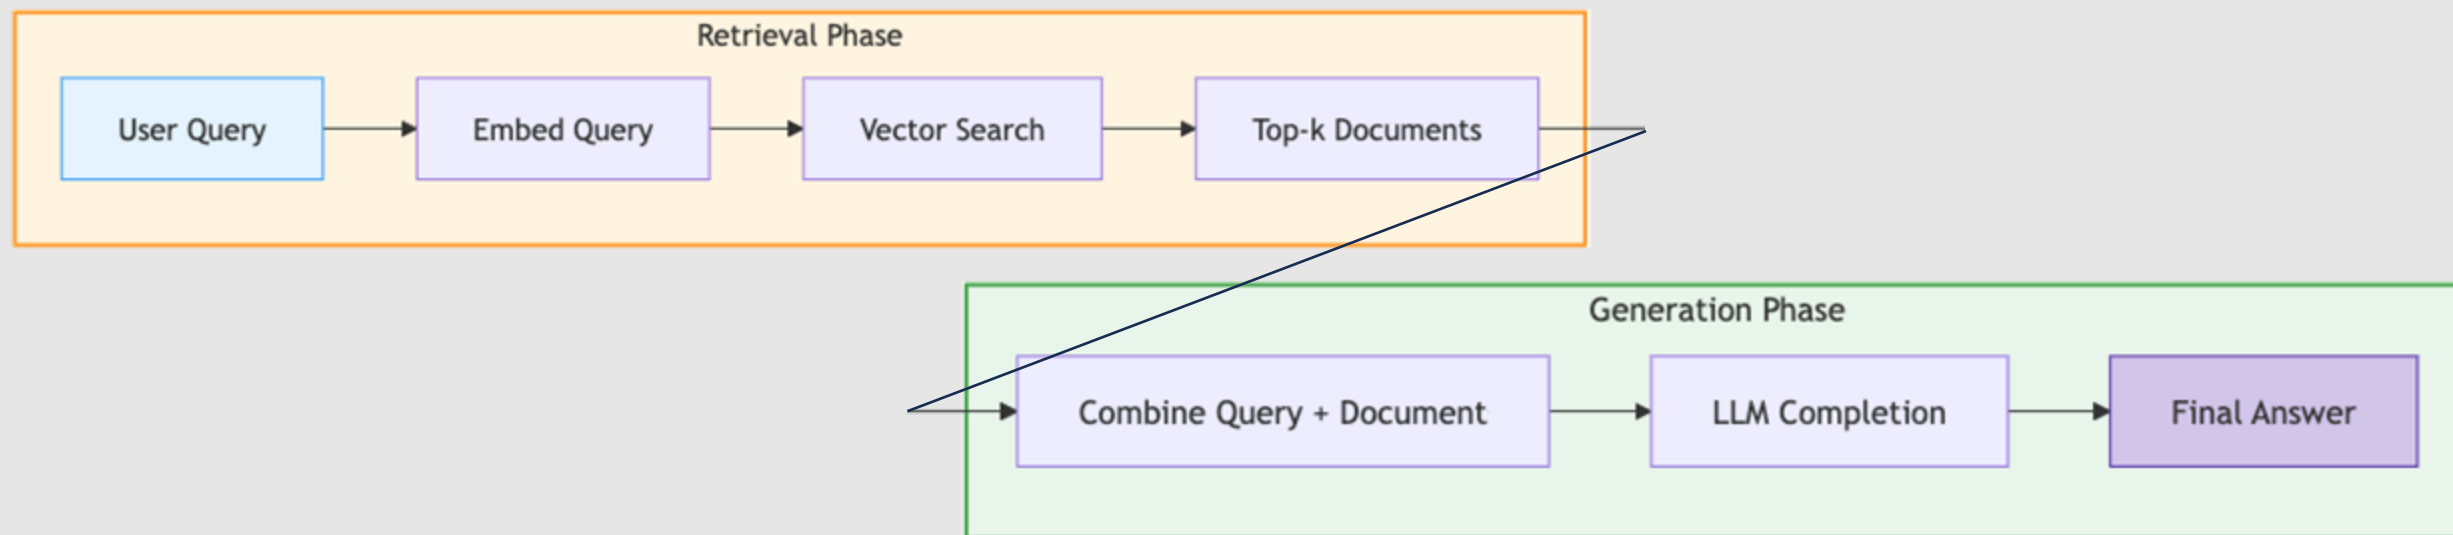
RAG





# What is RAG and how does it work?

- **Combination of retrieval and generation:** RAG combines information retrieval (like a search engine) with text generation (like a language model).
- **Uses external knowledge:** Instead of relying solely on pre-trained knowledge, RAG retrieves relevant documents or data from an external source (like a database or private knowledge bases) to generate more accurate and up-to-date responses.
- **Improves factual accuracy:** By pulling in real-time data or documents, RAG reduces the risk of generating factually incorrect or outdated information.
- **Two-step process:**
  - **Retrieve:** The model searches for relevant information from a knowledge source.
  - **Generate:** It then uses the retrieved data to create a coherent, contextually accurate answer.



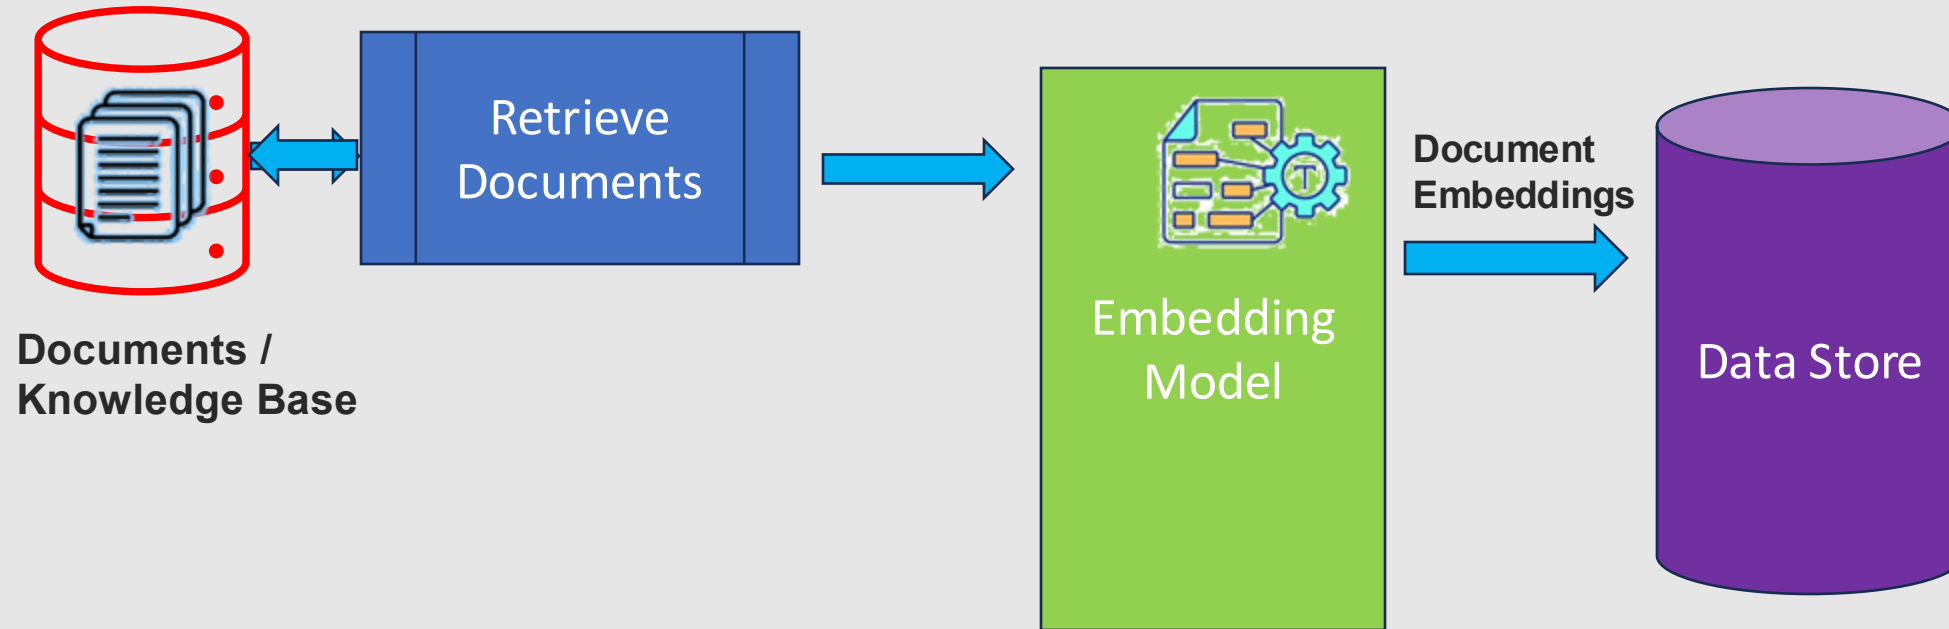
Source: <https://blogs.nvidia.com/blog/what-is-retrieval-augmented-generation/>



# How is RAG setup?

30

## *Doc Ingestion and Retrieval*



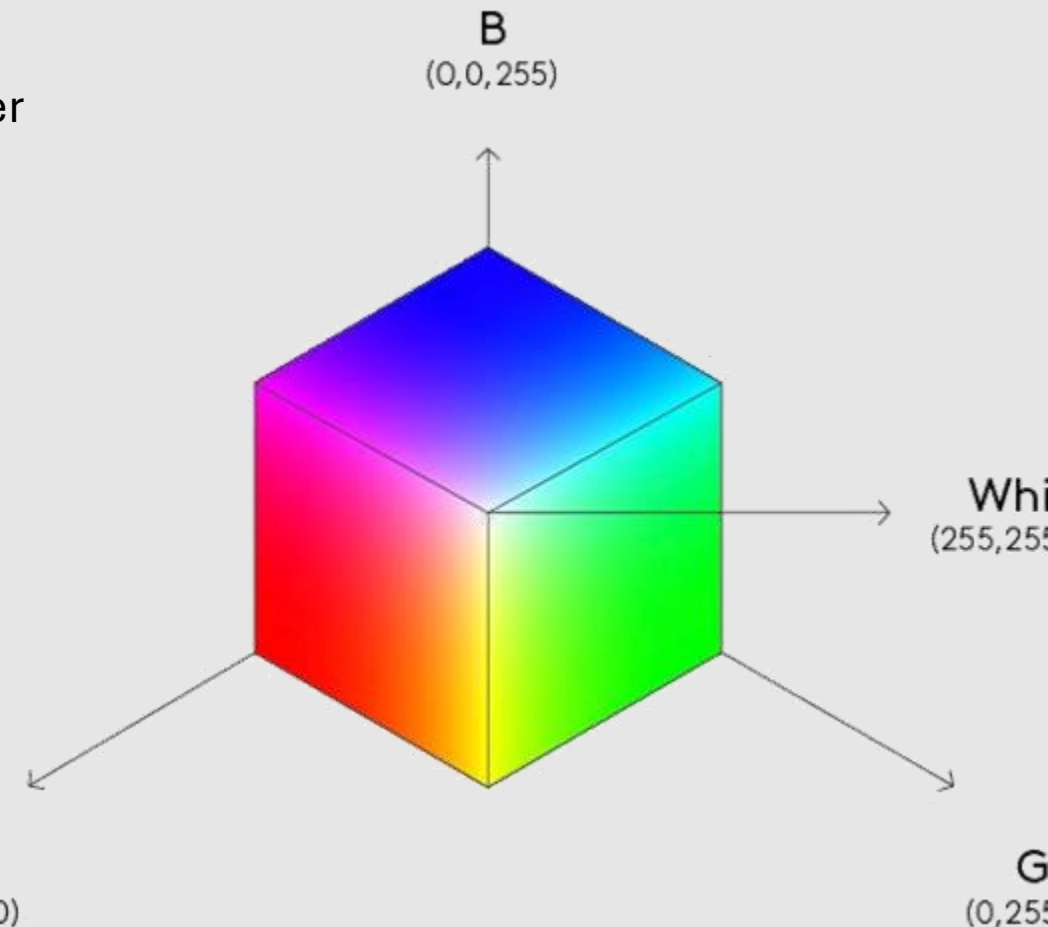
- You provide data sources and point application to them
- Info is retrieved from the data sources and tokenized, embedded and stored in a data store
- For queries/prompts, application gathers results (most relevant ones) from the vector database with your data



# Embeddings

31

- Embeddings represent text as sets of numeric data - tensors (lots of dimensions)
- Each dimension stores some info about the text's meaning, context, or syntactical aspects
- Words or sentences with similar meanings are stored closer together in the vector space
  - If two pieces of text are similar syntactically, they will have similar embeddings (smaller distance between their vectors)
- During training, models learn to place text with similar meanings closer together in the embedding space
- Common pre-trained models used for generating embeddings include BERT and variants (RoBERTa, DistilBERT)
- Once you have embeddings, you can use them for NLP tasks like semantic search, text classification, sentiment analysis

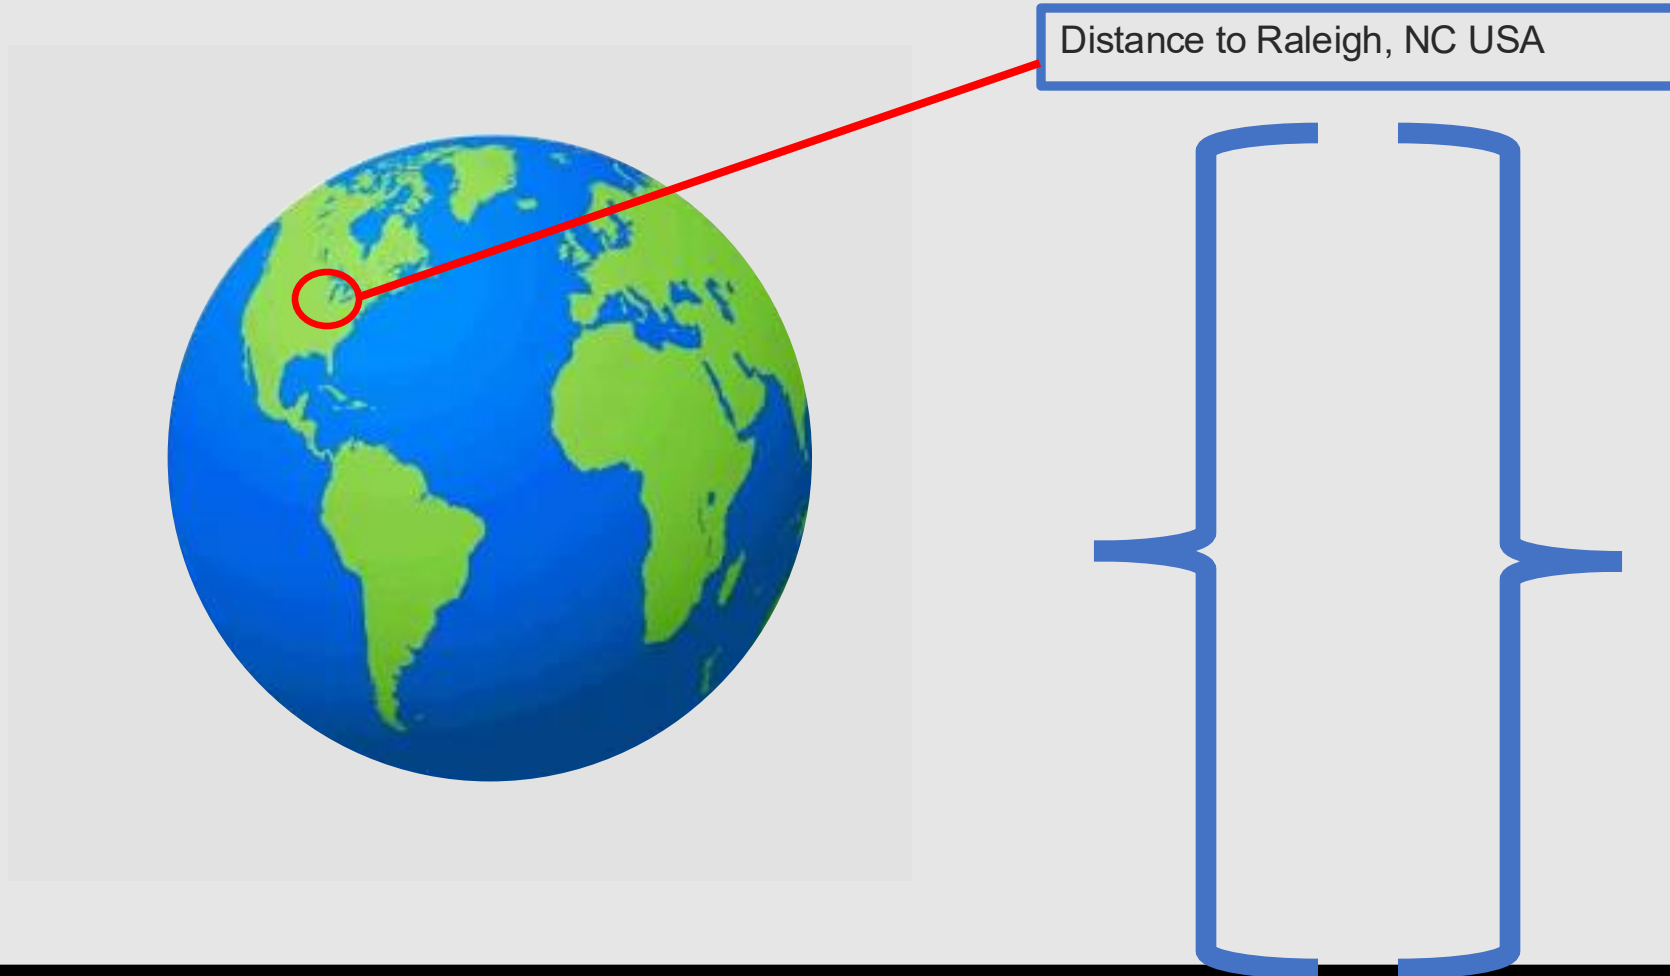






# Understanding vectors in AI

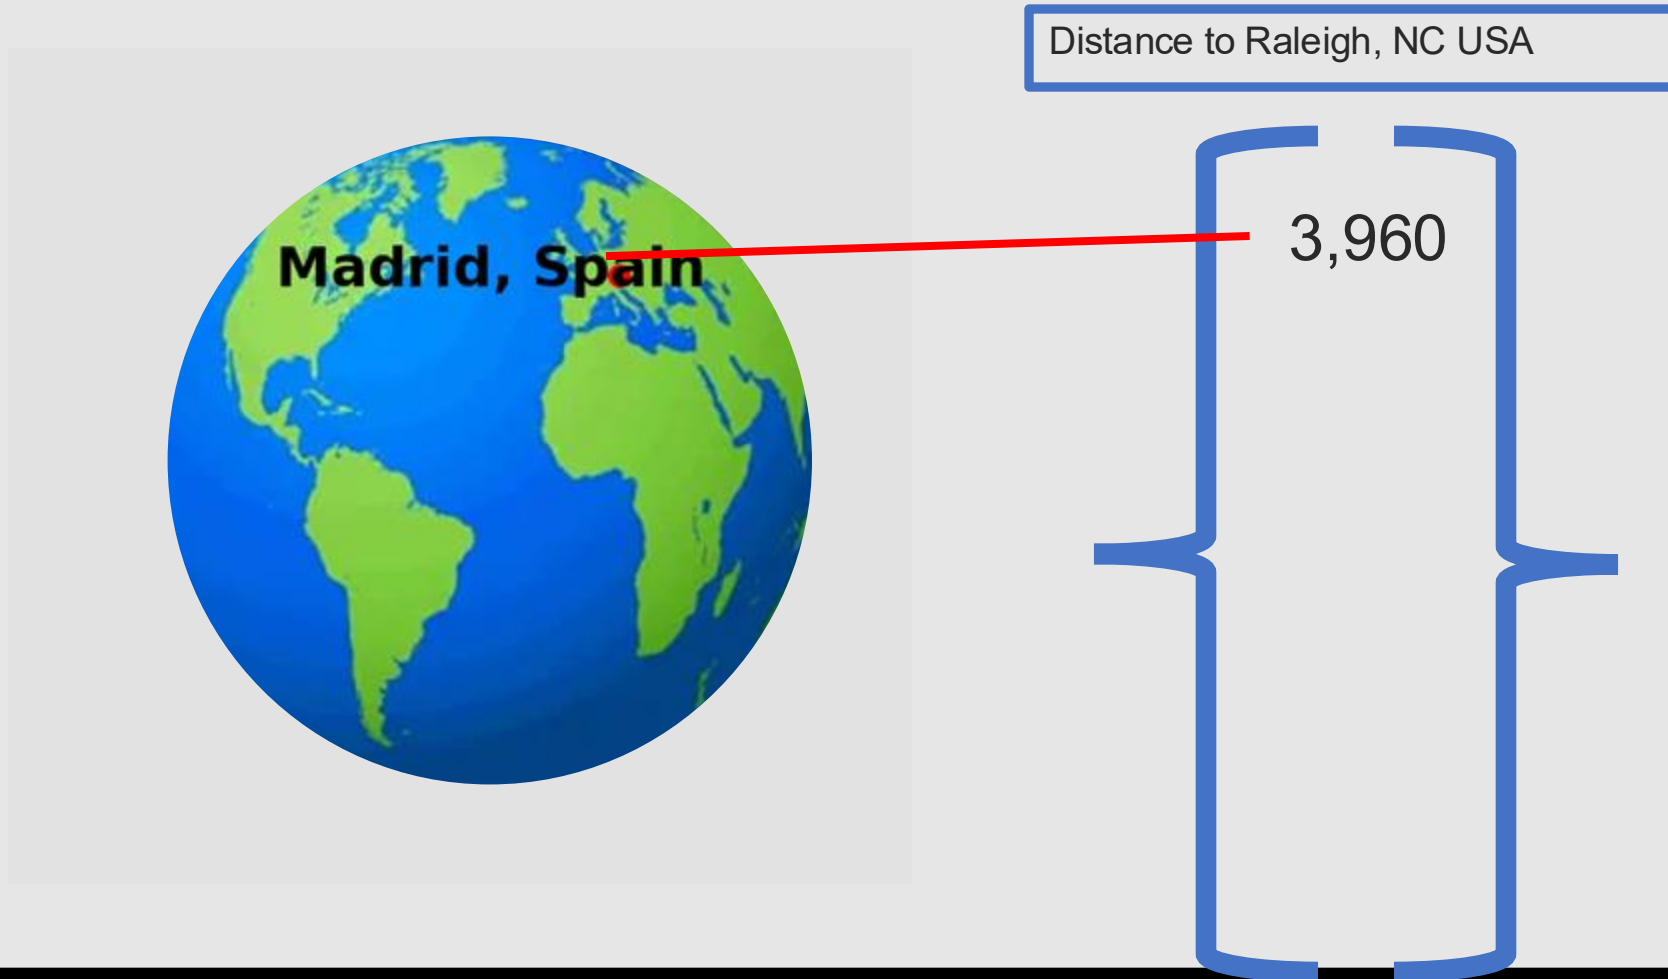
- Collection of data points that encapsulate an item's relationship to other items





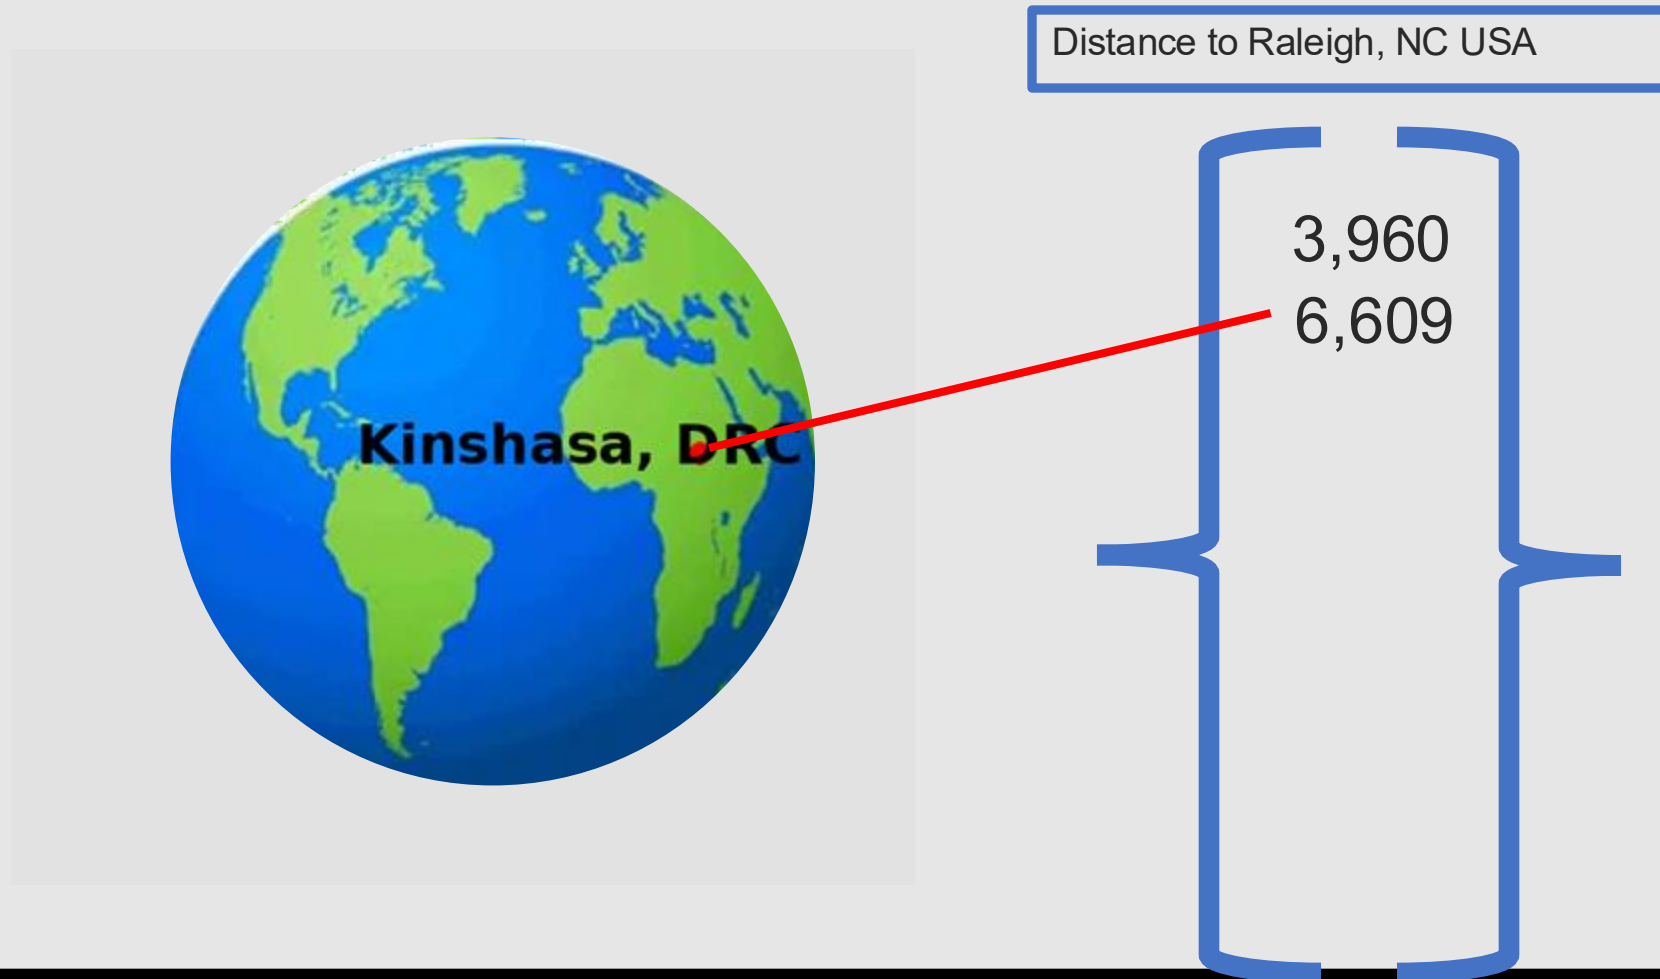
# Understanding vectors in AI

- Collection of data points that encapsulate an item's relationship to other items



# Understanding vectors in AI

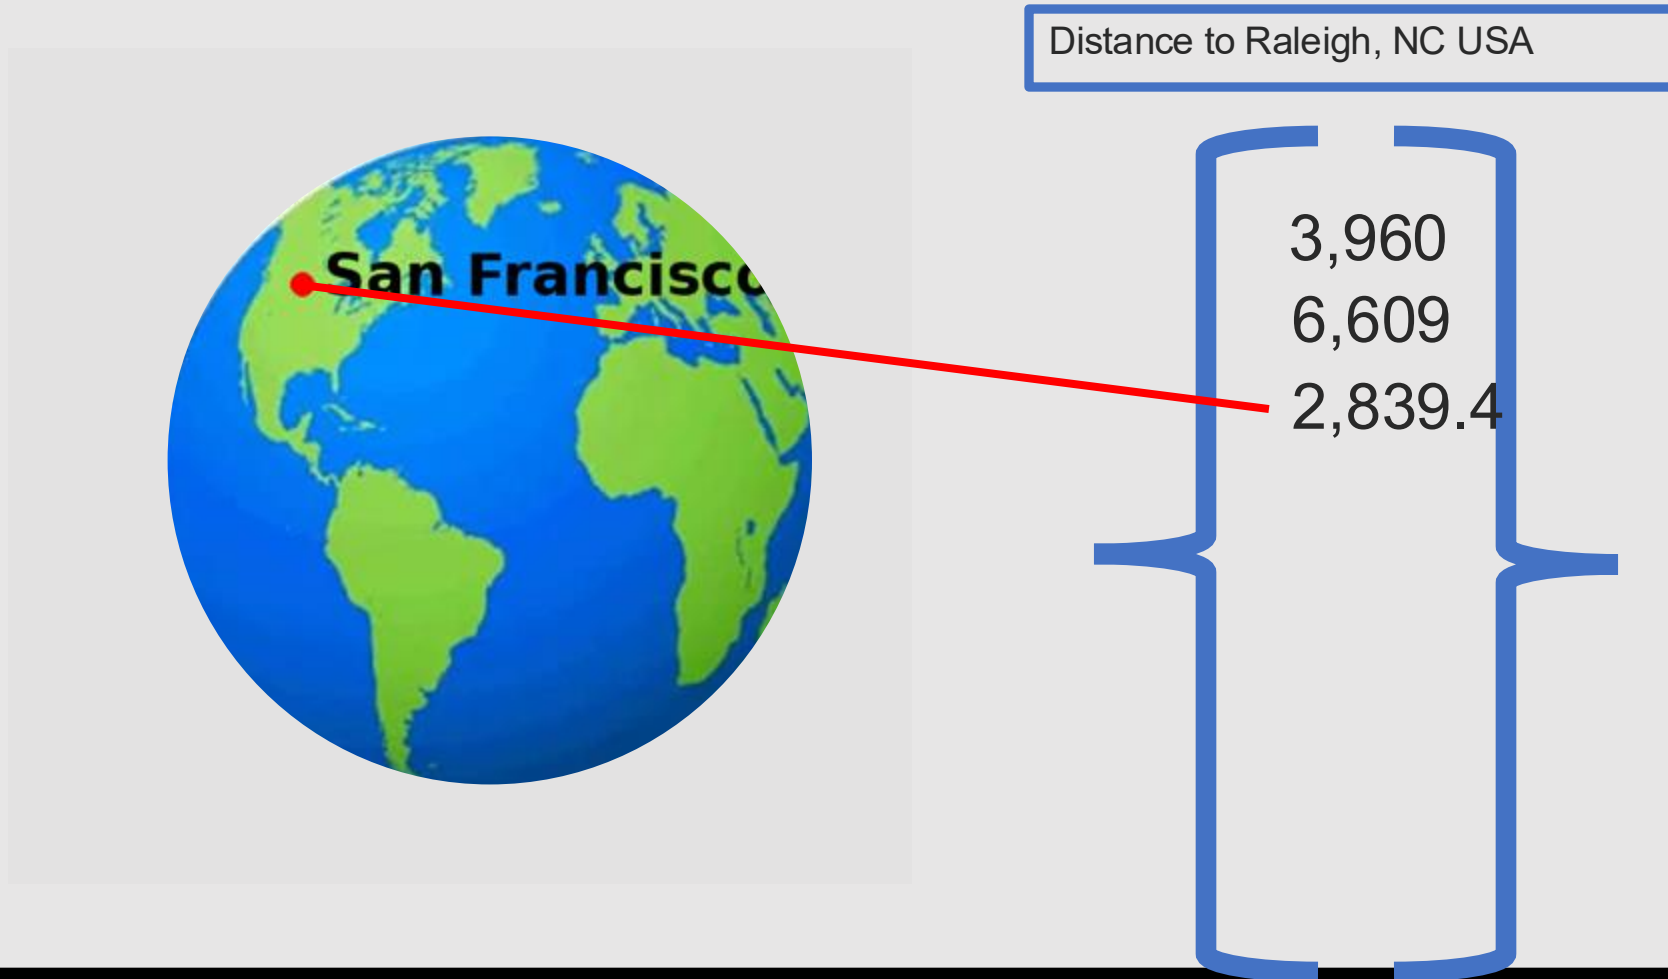
- Collection of data points that encapsulate an item's relationship to other items





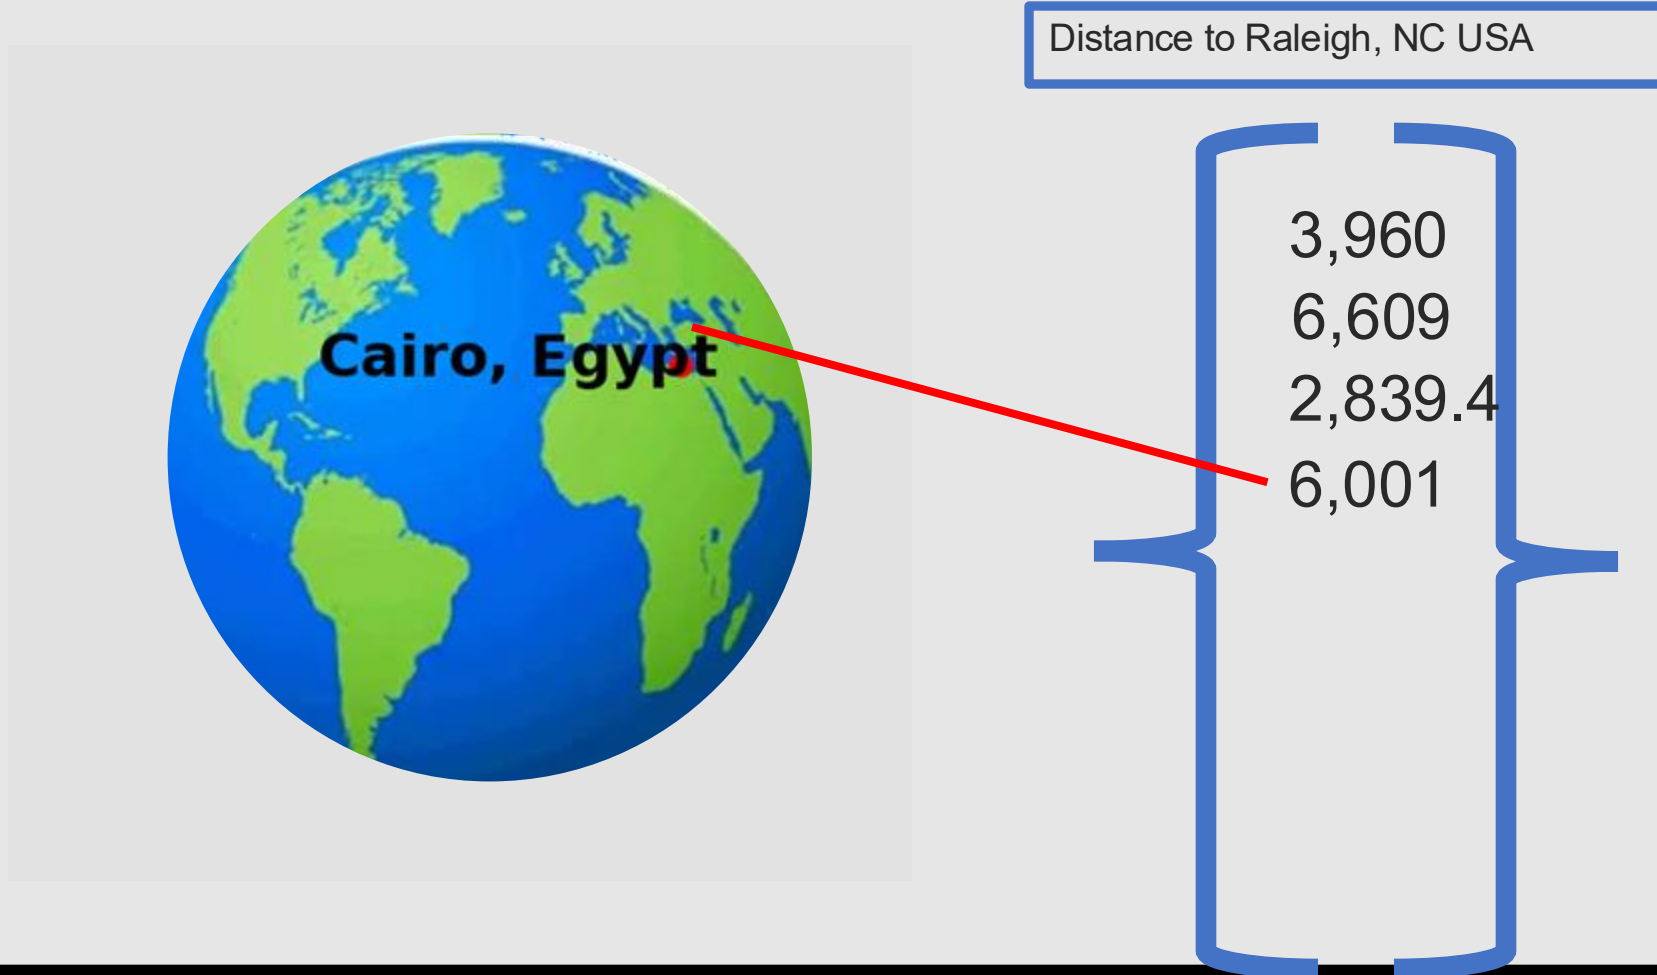
# Understanding vectors in AI

- Collection of data points that encapsulate an item's relationship to other items



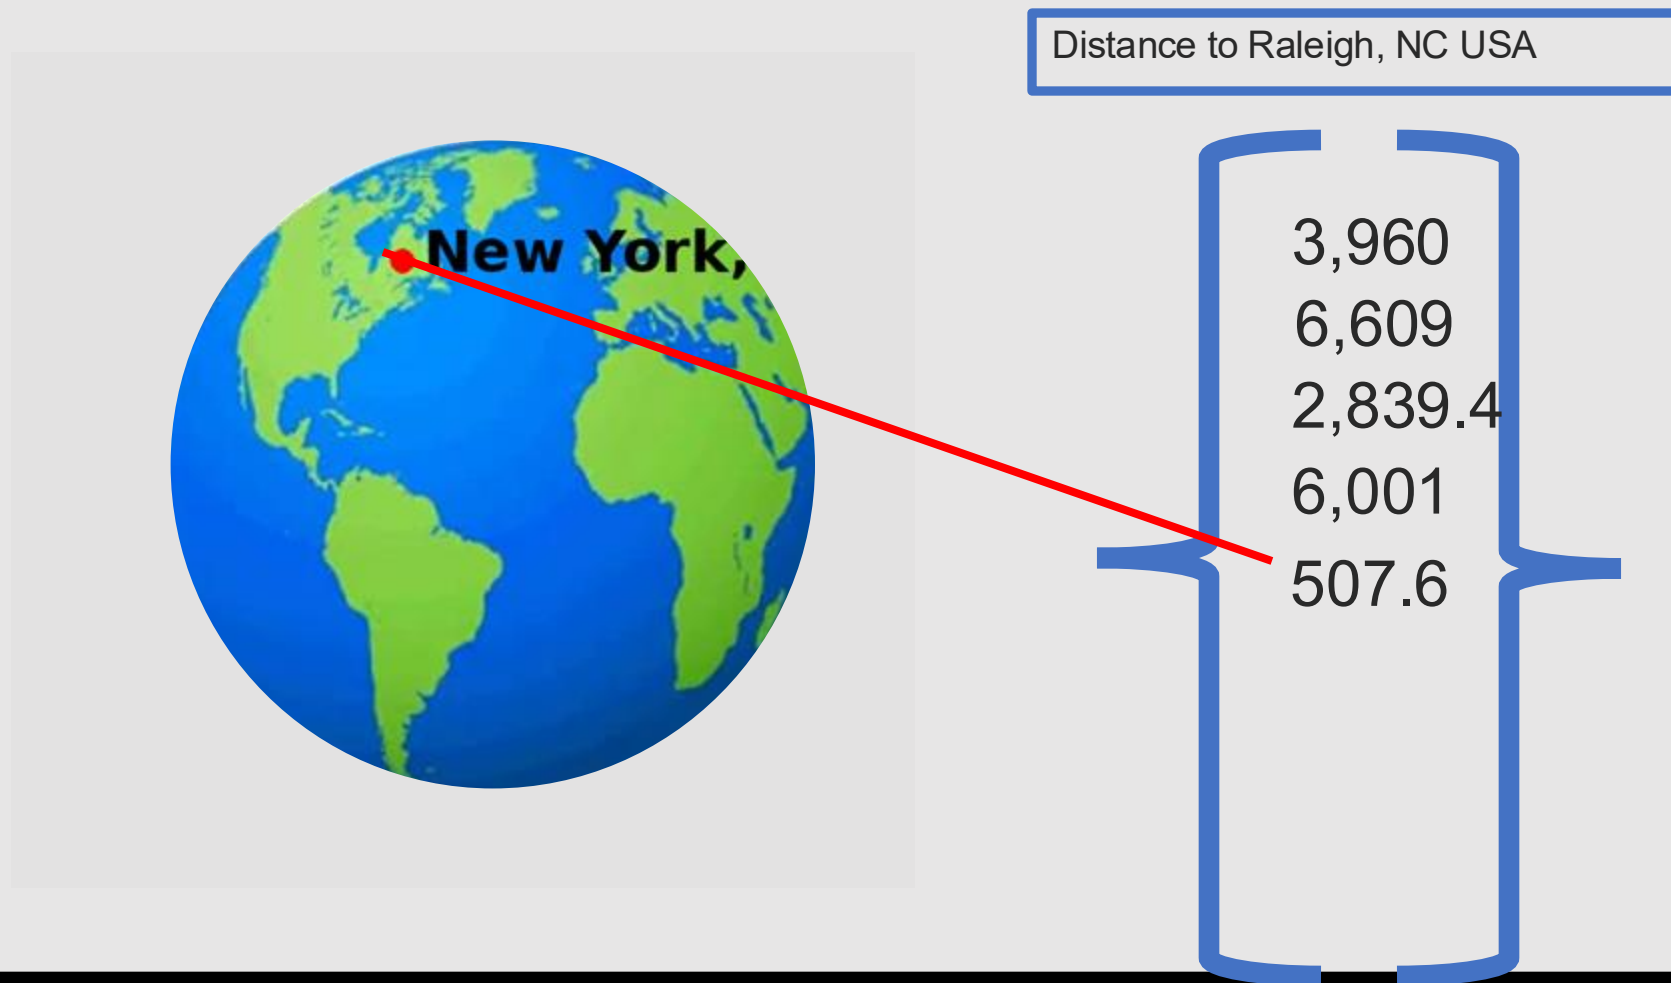
# Understanding vectors in AI

- Collection of data points that encapsulate an item's relationship to other items



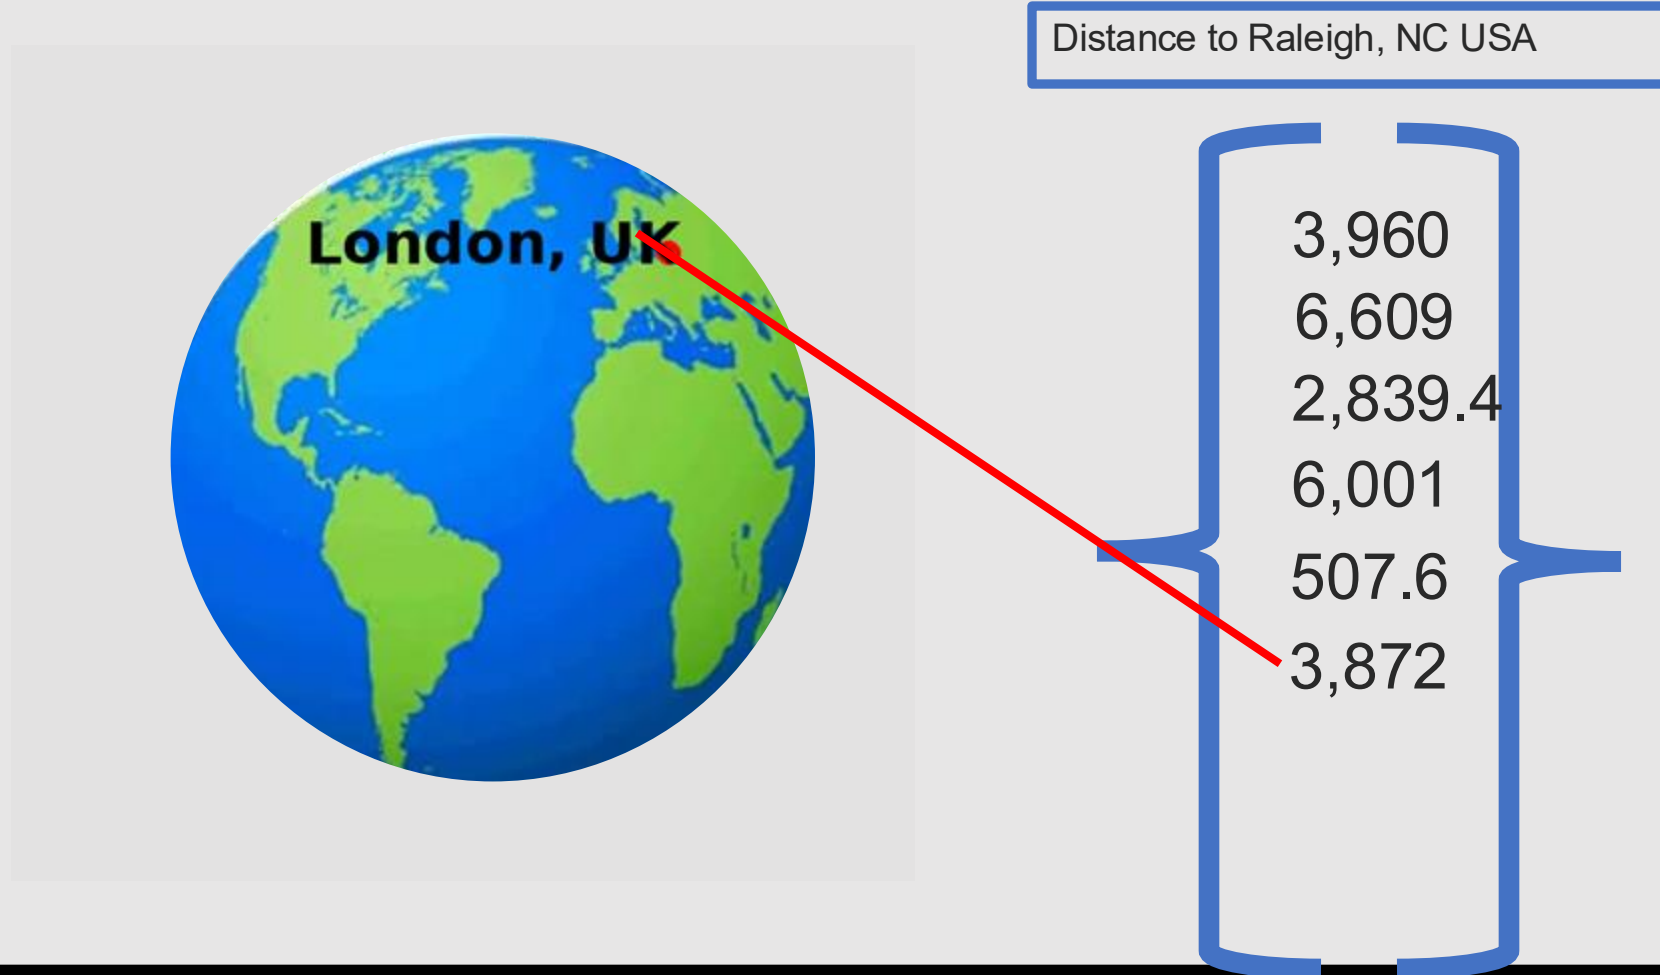
# Understanding vectors in AI

- Collection of data points that encapsulate an item's relationship to other items



# Understanding vectors in AI

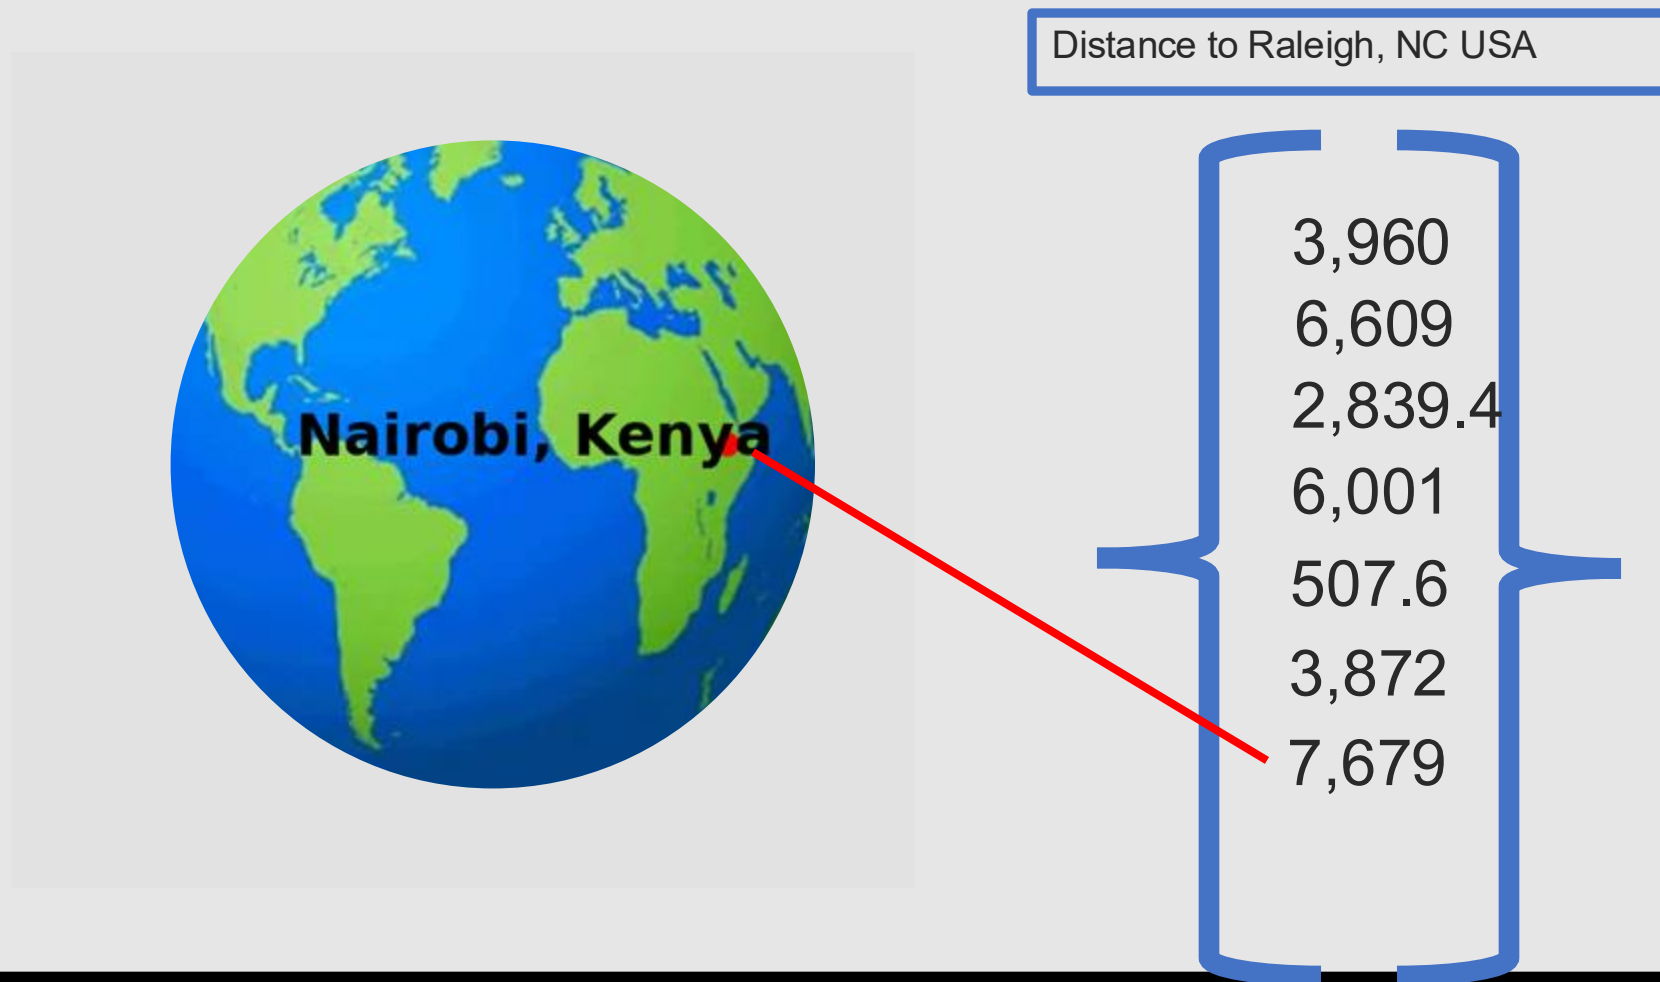
- Collection of data points that encapsulate an item's relationship to other items





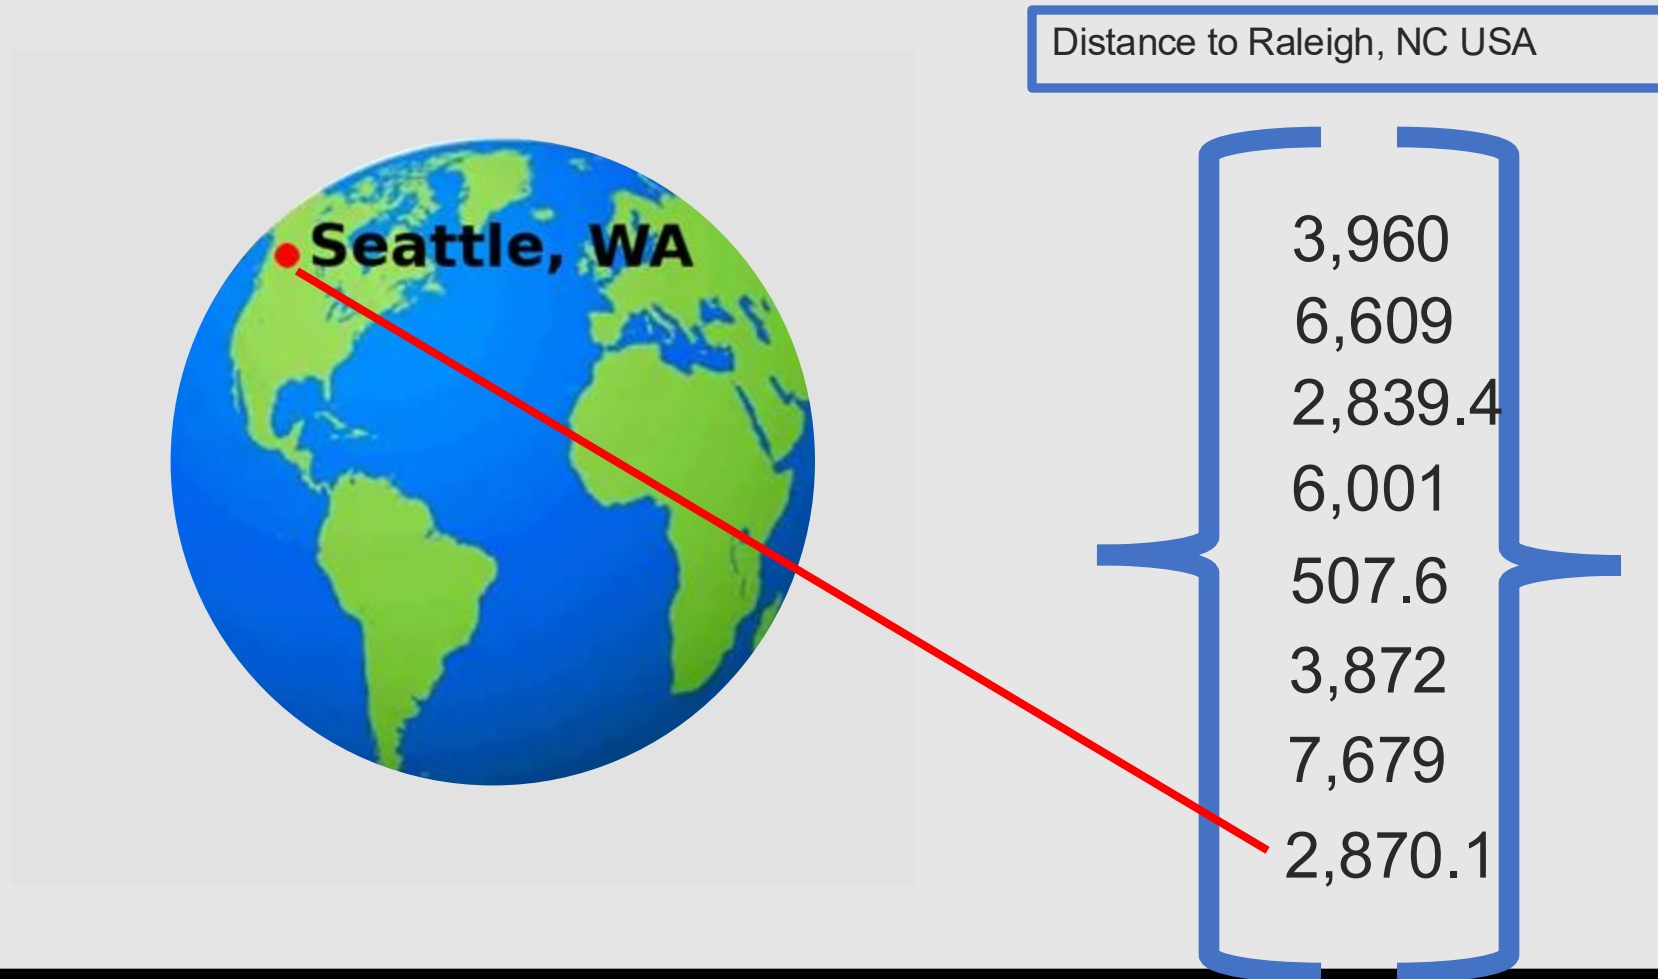
# Understanding vectors in AI

- Collection of data points that encapsulate an item's relationship to other items



# Understanding vectors in AI

- Collection of data points that encapsulate an item's relationship to other items





# Semantic meaning / relationships

- Suppose we have 3 words
- King and Queen are more similar to each other than they are to lunch
- In order for neural net to understand the relationships, each word needs to be represented as a vector
- Suppose each word is represented by a 2-dimensional vector

King  $\begin{bmatrix} -130.16 \\ 89.5 \end{bmatrix}$

Queen  $\begin{bmatrix} -115.43 \\ 95.2 \end{bmatrix}$

Lunch  $\begin{bmatrix} -89.5 \\ 34.3 \end{bmatrix}$




# Embedding space

- Plotting in 2-dimensional embedding space shows relationships
- Way to let NN understand relationships between words
- We want the NN to learn that King and Queen are more similar to each other than they are to lunch

King  $\begin{bmatrix} -130.16 \\ 89.5 \end{bmatrix}$

Queen  $\begin{bmatrix} -115.43 \\ 95.2 \end{bmatrix}$


Lunch  $\begin{bmatrix} -89.5 \\ 34.3 \end{bmatrix}$






# Searching for Vectors - similarity metrics


- 3 metrics commonly used to determine similarity of two vectors (2-dimensional representation)



**Cosine similarity** - measure the angle between two vectors; values from -1 to 1; 1 = both point in same direction; -1 point in opposite directions; 0 = orthogonal (perpendicular)




**Dot product / inner product** - measures how well 2 vectors align with each other; values from  $-\infty$  to  $\infty$ ; positive values indicate vectors are in same direction; negative values indicate opposite directions; 0 = orthogonal



**Euclidean distance** - measures the distance between two vectors; values from 0 to  $\infty$ ; 0 = identical; larger numbers farther apart

imagine 3 vectors -  $a, b, c$

$$a = \begin{bmatrix} .01 \\ .07 \\ .1 \end{bmatrix} \quad b = \begin{bmatrix} .01 \\ .08 \\ .11 \end{bmatrix} \quad c = \begin{bmatrix} .91 \\ .57 \\ .6 \end{bmatrix}$$



**Cosine similarity**

$$\text{sim}(u, v) = \frac{u \cdot v}{\|u\| \|v\|} = \frac{\sum_{i=1}^n a_i b_i}{\sqrt{\sum_{i=1}^n u_i^2} \sqrt{\sum_{i=1}^n v_i^2}}$$

0.0141

$$\begin{aligned} \text{sim}(a, b) &= \frac{(a_1 * b_1) + (a_2 * b_2) + (a_3 * b_3)}{\sqrt{a_1^2 + a_2^2 + a_3^2} \sqrt{b_1^2 + b_2^2 + b_3^2}} \\ &= \frac{(0.01 * 0.01) + (0.07 * 0.08) + (0.1 * 0.11)}{\sqrt{0.01^2 + 0.07^2 + 0.1^2} \sqrt{0.01^2 + 0.08^2 + 0.11^2}} \end{aligned}$$

**Dot product / inner product**

$$u \cdot v = |u| |v| \cos \theta = \sum_{i=1}^n a_i b_i$$

0.0167

Dot product formula

$$\begin{aligned} a \cdot b &= (a_1 b_1) + (a_2 b_2) + (a_3 b_3) \\ &= (0.01 * 0.01) + (0.07 * 0.08) + (0.1 * 0.11) \end{aligned}$$

**Euclidean distance**

$$d(u, v) = \sqrt{\sum_{i=1}^n (u_i - v_i)^2}$$

0.9998

Euclidean distance formula



$$\begin{aligned} d(a, b) &= \sqrt{(b_1 - a_1)^2 + (b_2 - a_2)^2 + (b_3 - a_3)^2} \\ &= \sqrt{(0.01 - 0.01)^2 + (0.08 - 0.07)^2 + (0.11 - 0.1)^2} \end{aligned}$$

edit: <https://towardsdatascience.com/similarity-metrics-in-nlp-acc0777e234c>





# Visualizing Embeddings and Vector Similarity



source: [https://projector.tensorflow.org/?config=https://gist.githubusercontent.com/martin-labrecque/4483ff5a104f0b56417585c3bc9a12f1/raw/57348e12a70c8d70c2c573d3dbc0122ac077556b/journaux\\_config.json](https://projector.tensorflow.org/?config=https://gist.githubusercontent.com/martin-labrecque/4483ff5a104f0b56417585c3bc9a12f1/raw/57348e12a70c8d70c2c573d3dbc0122ac077556b/journaux_config.json)













# Vectors and relationships example

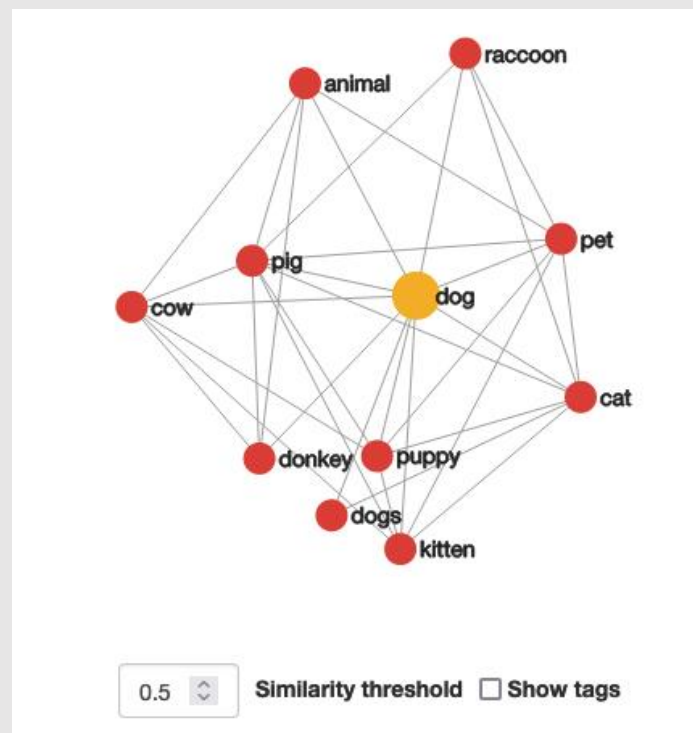
- Query - what words are related to "dog" in model "English Wikipedia"?

Word frequency

☒ High ☒ Medium ☐ Low

1. **puppy** NOUN 0.6909 
2. **cat** NOUN 0.6893 
3. **pet** NOUN 0.6779 
4. **pig** NOUN 0.6680 
5. **dogs** NOUN 0.6174
6. **animal** NOUN 0.6066 
7. **donkey** NOUN 0.6060 
8. **raccoon** NOUN 0.6021 
9. **cow** NOUN 0.5953
10. **kitten** NOUN 0.5947 

• We show only the associates of the same part of speech as your query. All associates can be found at the [Similar Words](#) tab.



Show the raw vector of «dog» in model

MOD\_enwiki\_upos\_skipgram\_300\_2\_2021:

[-0.03301828354597092, 0.05134638026356697, 0.0036009703762829304, -0.04066073149442673, 0.10361430048942566, 0.013021323829889297, 0.028161464259028435, -0.0027567853685468435, 0.03388035297393799, -0.044882044196128845, 0.005169689189642668, -0.05818631127476692, 0.0533536821603775, 0.016616210341453552, 0.02030780538916588, -0.008570297621190548, -0.10925538837909698, -0.0708925873041153, 0.04675082117319107, -0.03091960959136486, -0.05172094330191612, 0.04471702128648758, 0.008674593642354012, -0.01816382259130478, 0.05909318849444389, 0.10409023612737656, 0.05633684620261192, -0.024881813675165176, 0.01872968301177025, 0.007228093687444925, -0.023127363994717598, 0.01528552919626236, -0.0643191784620285, -0.010359424166381359, -0.06104437634348869, -0.13868044316768646, -0.023004498332738876, 0.0038427673280239105, -0.021551262587308884, -0.03467748314142227, 0.010687021538615227, -0.017304275184869766, 0.026886526495218277, -0.0030398862436413765, -0.03685504570603371, -0.06017328053712845, 0.047442398965358734, -0.10714898258447647, 0.14808930456638336, -0.06579480320215225, -0.004342162515968084, 0.06226382404565811, 0.08031187951564789, -0.055930640548467636, -0.07030591368675232, 0.015474628657102585, 0.05367768555879593, 0.0917837843298912, 0.031899698078632355, 0.055091146379709244, -0.025078952312469482, -0.048126623034477234, -0.09730836749076843, -0.07128141075372696, 0.019415033981204033, -0.025872433558106422, -0.01761292852461338, 0.015608762390911579, -0.029876720160245895, -0.008602319285273552, 0.049825914204120636, 0.06784739345312119, 0.005586292129009962, -0.07148509472608566, -0.03097137063741684, -0.020296750590205193, 0.05099814385175705, 0.14920306205749512, 0.03855258598922683, -0.0818730816245079, -0.06150494143366814]

Source: [http://vectors.nlp.eu/explore/embeddings/en/MOD\\_enwiki\\_upos\\_skipgram\\_300\\_2\\_2021/dog\\_NOUN/](http://vectors.nlp.eu/explore/embeddings/en/MOD_enwiki_upos_skipgram_300_2_2021/dog_NOUN/)



# Vector Databases







# Vector Databases

47

- Specialized database that index and stores *vector embeddings*
- Useful for
  - fast retrieval
  - similarity search
- Offer comprehensive data management capabilities
  - metadata storage
  - filtering
  - dynamic querying based on associate metadata
- Scalable and can handle large volumes of vector data
- Support real-time updates
- Play key role in AI and ML applications

Vector Database



Weaviate



milvus



drant



Pinecone



SingleStore



Chroma



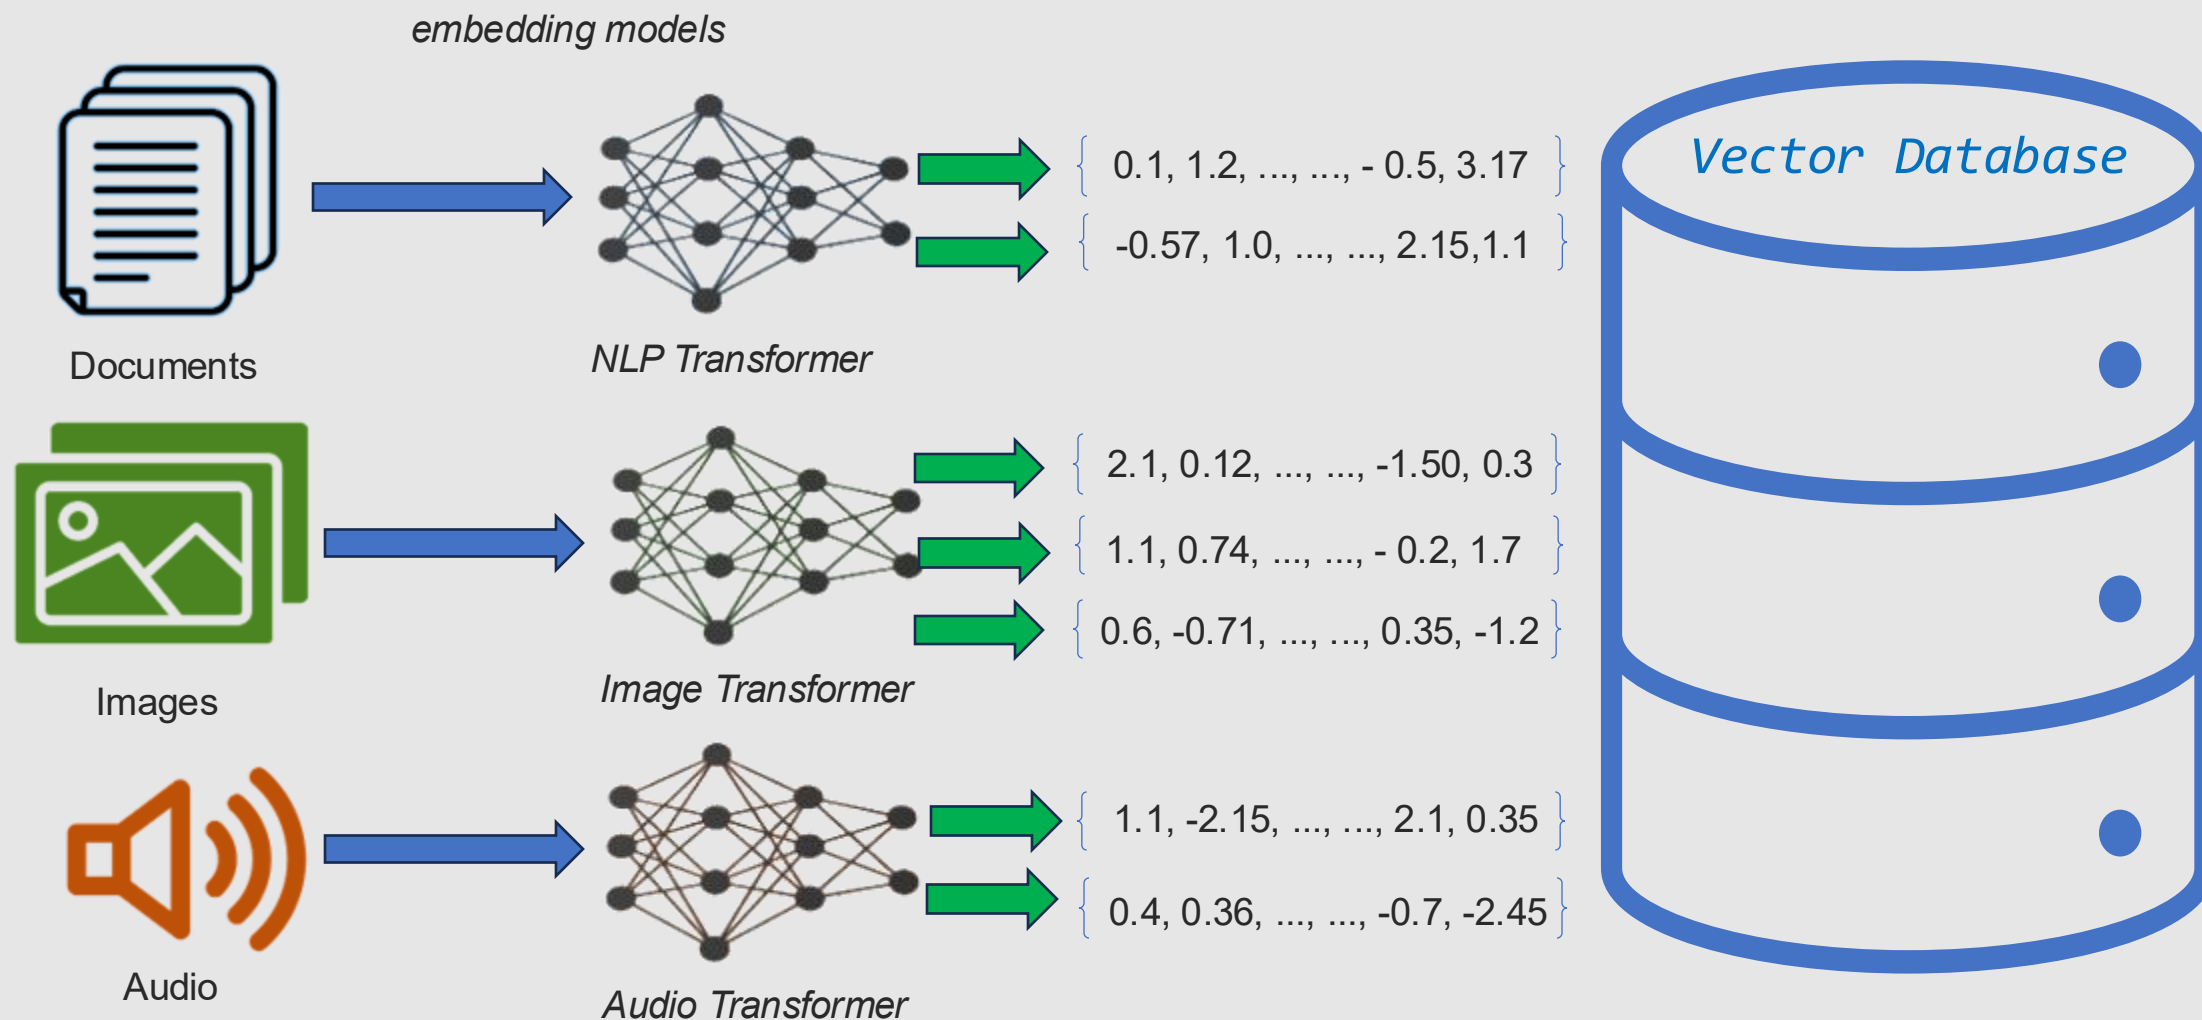
zilliz



# How data gets into Vector Databases

48

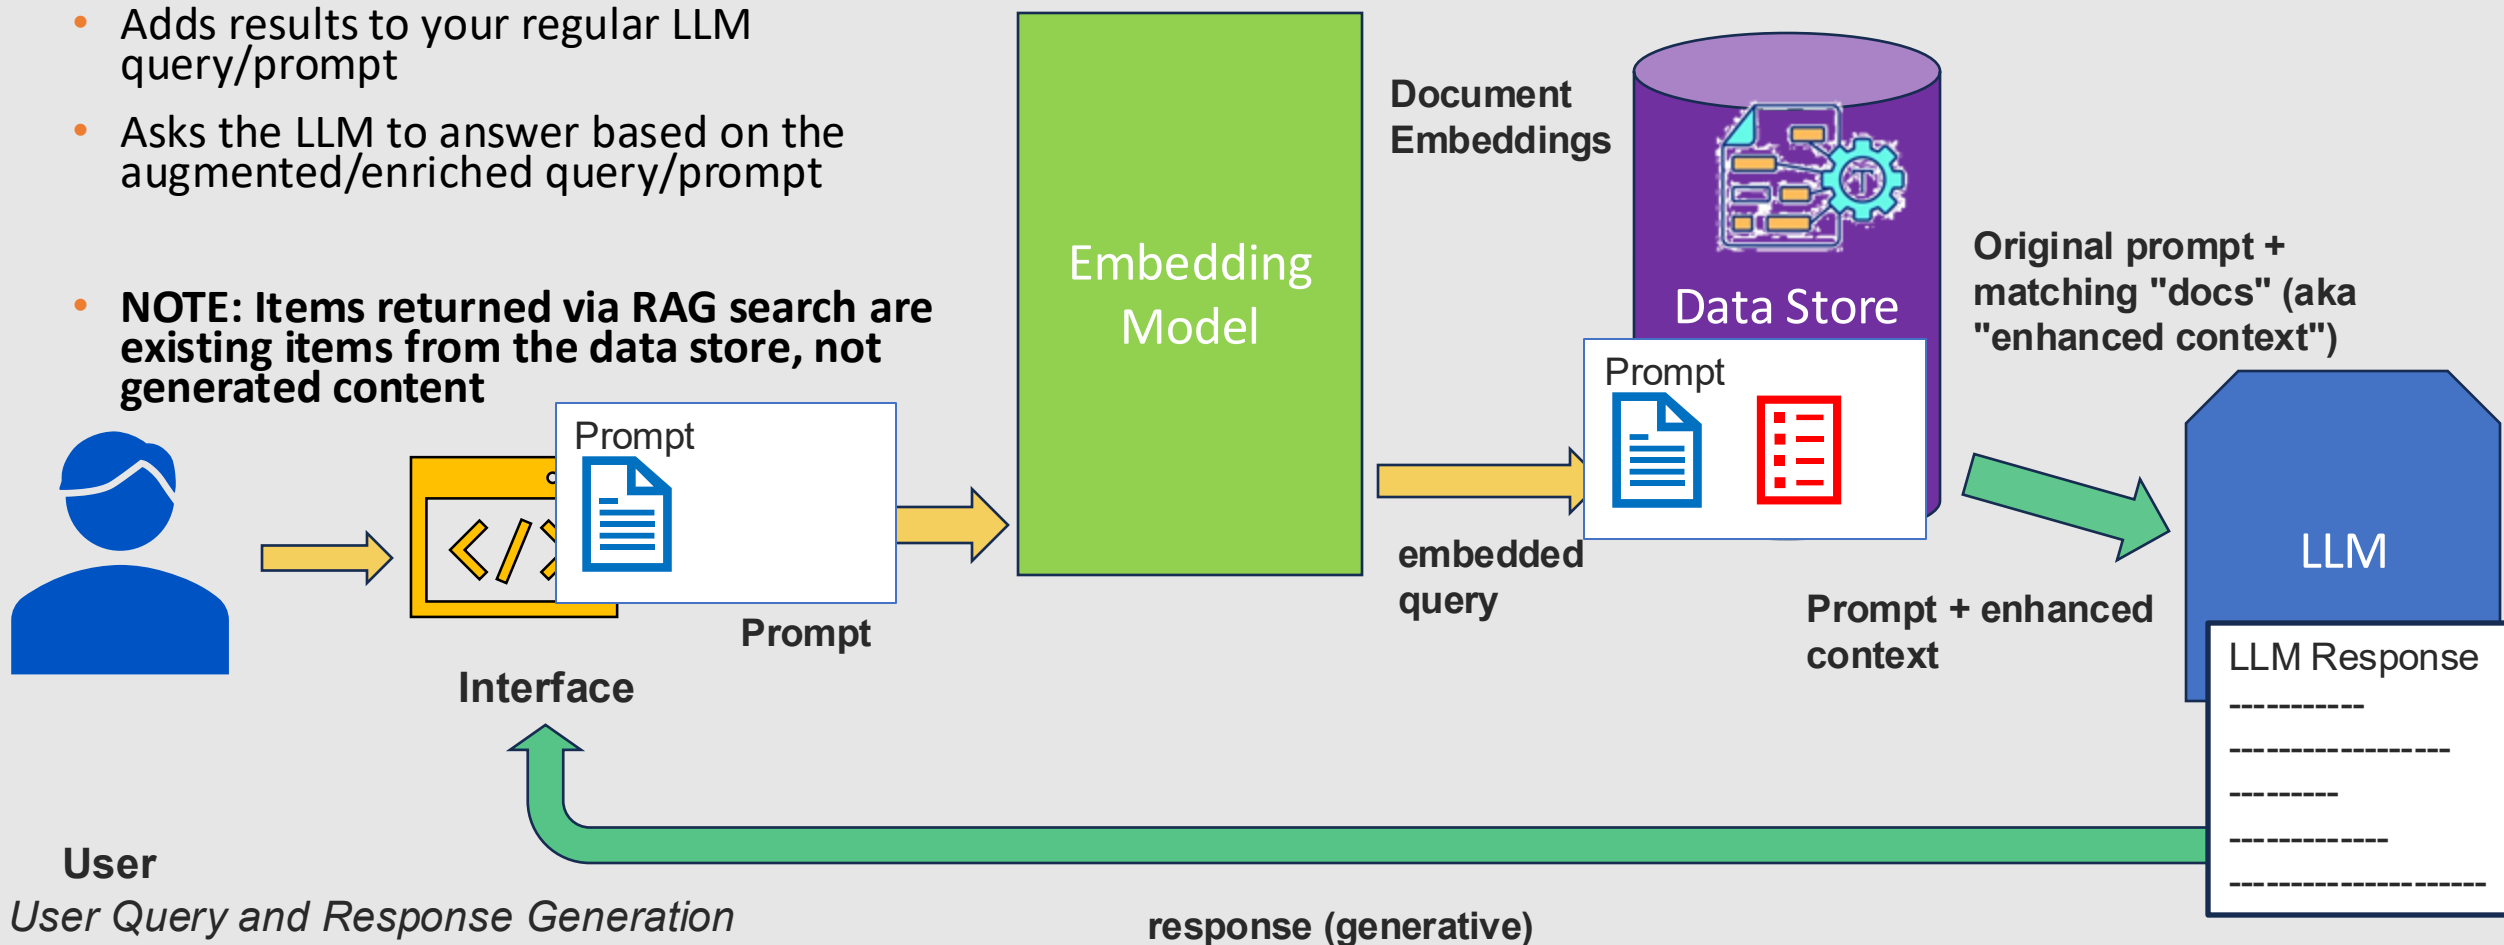
- Data is input, converted to embeddings (vectors) and stored
- Queries are input, converted to embeddings (vectors) and then **similarity metrics** are used to find results ("nearest neighbors")





# How does RAG work?

- For queries/prompts, application gathers results (most relevant ones) from the vector database with your data
- Adds results to your regular LLM query/prompt
- Asks the LLM to answer based on the augmented/enriched query/prompt
- **NOTE: Items returned via RAG search are existing items from the data store, not generated content**





## Demo #3 – Adding RAG to our code

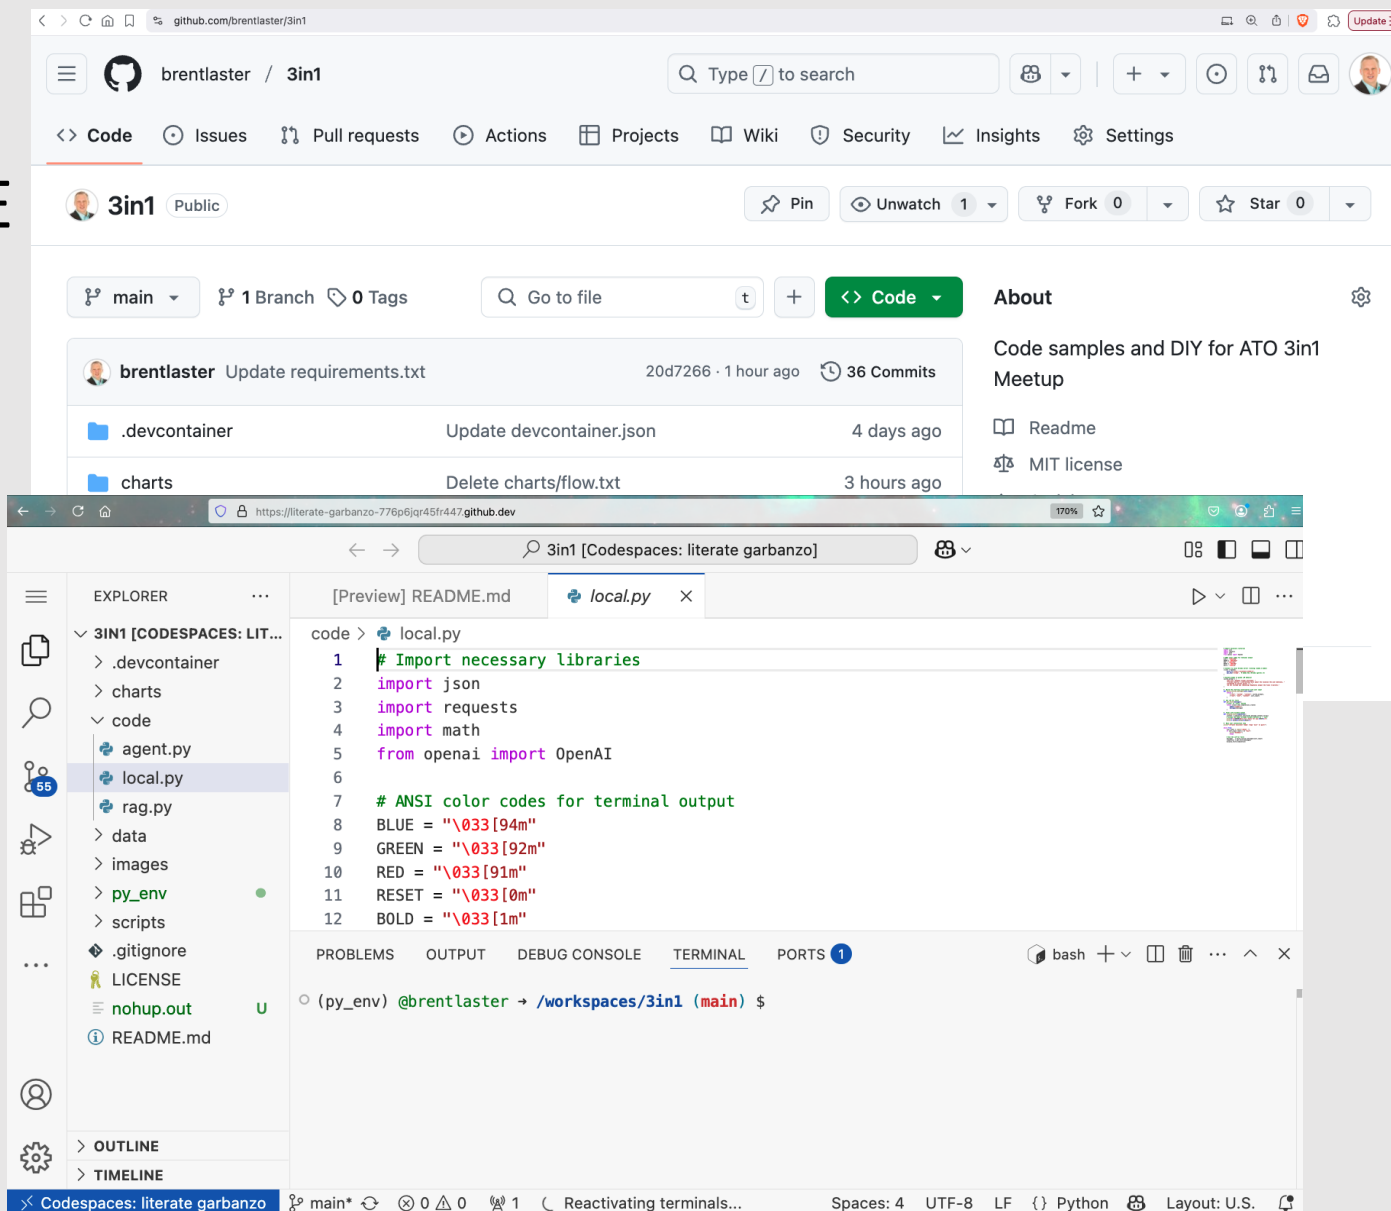
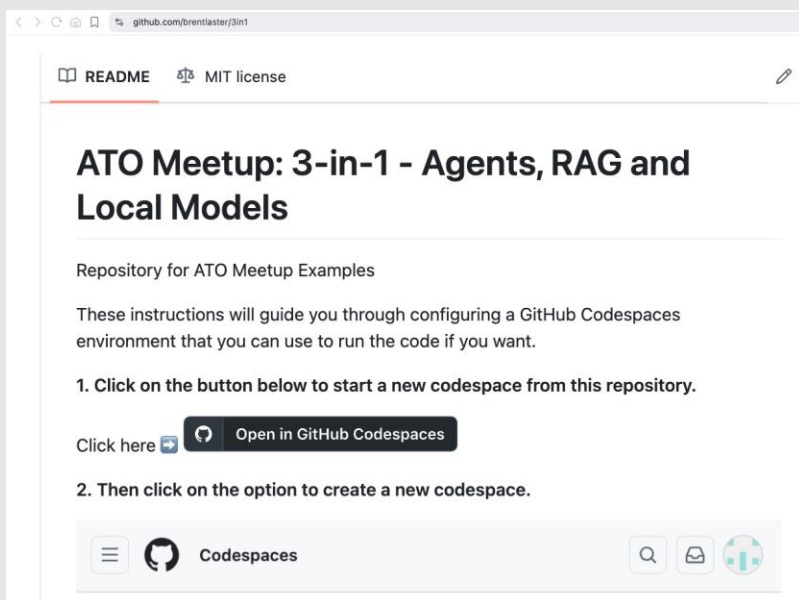




# DIY – github.com/brentlaster/3in1

51

- Fork if desired
- Click on button in README to start codespace
- Follow guide.md





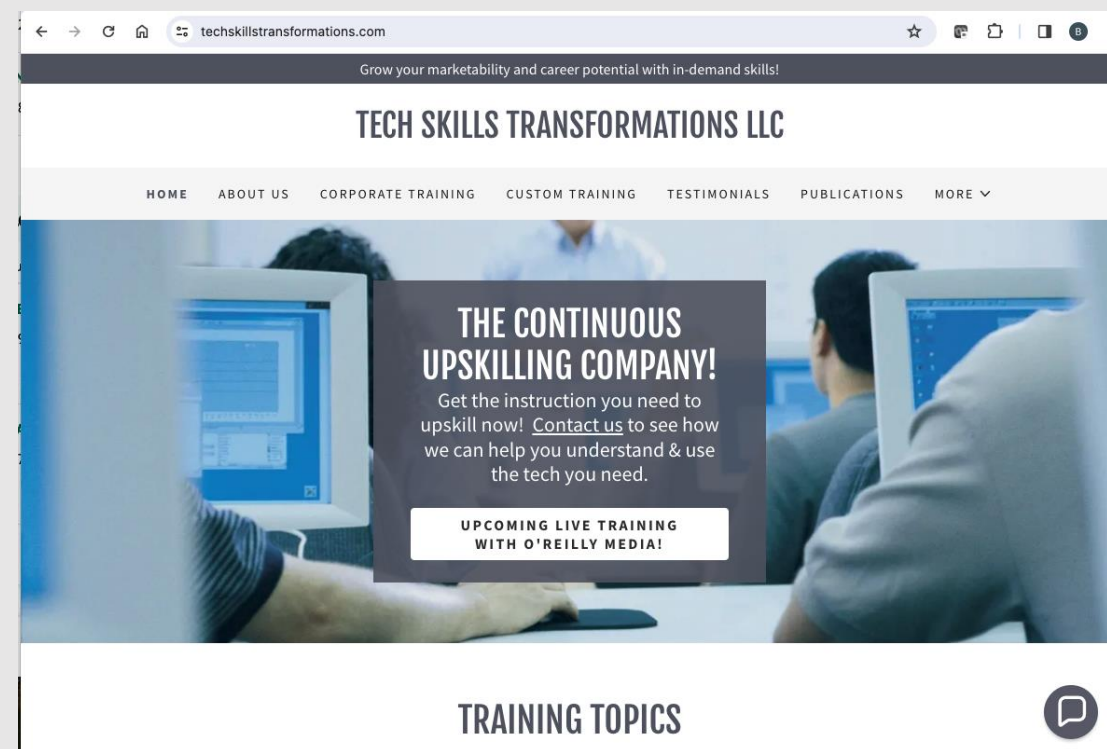







# That's all - thanks!

52

**Contact:** [training@getskillsnow.com](mailto:training@getskillsnow.com)

[techskillstransformations.com](https://techskillstransformations.com)  
[getskillsnow.com](https://getskillsnow.com)



❑ LinkedIn: [brentlaster](#)  
❑ X: [@BrentCLaster](#)  
❑ Bluesky: [brentclaster.bsky.social](#)  
❑ GitHub: [brentlaster](#)

